Overview of the package LMMstar

Brice Ozenne

May 12, 2024

This vignette describes the main functionalities of the **LMMstar** package. This package implements specific types of linear mixed models, mainly useful when having repeated observations over a discrete variable: $\boldsymbol{Y} = (Y_1, \ldots, Y_T)$ where T can be for example be time (chronological ordering of the repetitions) or brain region (arbitrary ordering of the repetitions). Denoting by $\boldsymbol{X}$ the associated covariates and $\boldsymbol{\varepsilon} = (\varepsilon_1, \ldots, \varepsilon_T)$, the model can be written:

$$\boldsymbol{Y} = \boldsymbol{X}\beta + \boldsymbol{\varepsilon} \text{ where } \varepsilon \sim \mathcal{N}(0, \Omega)$$

where β are the mean parameters and the residual variance-covariance matrix, Ω, depends on a set of variance-covariance parameters (say γ) distinct of β. Key assumptions are:

- we observe n independent replicates of $\mathcal{O} = (\boldsymbol{Y}, \boldsymbol{X})$, i.e. at the cluster level, observations $(\mathcal{O}_1, \ldots, \mathcal{O}_n)$ are independent. The replicates should also be identically distributed up to a categorical variable (called strata variable in the following).

- the residual variance is independent of the mean value.

Additional assumptions are necessary in presence of missing values, typically correct specification of the conditional mean to have consistent estimates of the mean parameters. This case will sometimes be examplified by considering that only last outcome may be missing: the conditional mean $\mathbb{E}[Y_T|Y_1, Y_2, \ldots, Y_{T-1}]$ is then abreviated as $\mathbb{E}[Y_T|Y_{T-1}]$. Note that we do not require the residuals to be normally distributed to have valid estimates or statistical inference in large samples.

To get start, one should load the **LMMstar** package in the ℝ session:

```
library(LMMstar)
```

This package is under active development. Newer package versions may include additional functionalities and fix previous bugs. The version of the package that is being used for this overview is:

```
utils::packageVersion("LMMstar")
```

[1] '1.1.0'

It is recommanded to also the following packages, as some of the methods implemented in the package are relative to a generic method implmented in other packages:

```
library(ggplot2) ## autoplot method
library(nlme) ## ranef method
library(lava) ## iid, information, manifest methods
```

The user interface of the **LMMstar** package is made of the following functions:

- functions to describe or visualize the dataset:

  - `scatterplot` to visualize the marginal and bivariate distribution of continuous variables.
  - `summarize` to compute summary statistics, possibly stratified on a categorical variable.
  - `summarizeNA` to identify missing data patterns.
  - `partialCor` to compute partial correlation between two variables.

- the function `mt.test` to perform multiple Student's t-Tests and adjust the results for multiple testing.

- the function `lmm` is the main function of the package which fits linear mixed models. The user can interact with *lmm* objects using:

  - `anova` to perform Wald tests, i.e. test linear combinations of coefficients ($\widehat{\beta}_1 + \widehat{\beta}_2 = 0$ or $\widehat{\beta}_1 = \widehat{\beta}_2 = 0$). The output obtained with different `lmm` can be combined using `rbind`.
  - `coef` to extract the estimated model parameters ($\widehat{\beta}$ and possibly $\widehat{\gamma}$).
  - `confint` to extract the estimates with their confidence intervals.
  - `effects` to evaluate marginal effects, e.g. $\mathbb{E}\left[\mathbb{E}\left[Y|X_1=1\right] - \mathbb{E}\left[Y|X_1=0\right]\right]$ when $\boldsymbol{X} = (X_1, X_2)$.
  - `estimate` to test non-linear combinations of coefficients (Wald test via a first order delta method, e.g. $\widehat{\beta}_1/\widehat{\beta}_2 = 1$).
  - `fitted` to output the fitted mean ($X\widehat{\beta}$) or the conditional mean for observations with missing outcome (e.g. $X\widehat{\beta} + \widehat{\mathbb{E}}[\varepsilon_T|\varepsilon_{-T}]$).
  - `iid` to extract the influence function of the estimated parameters (φ), which satisfies
    $\sqrt{n}(\widehat{\beta} - \beta) = \frac{1}{\sqrt{n}}\sum_{i=1}^{n}\varphi\left(\mathcal{O}_i\right) + o_p(1)$
  - `levels` to extract the reference level for the mean structure. (i.e. what (`Intercept`) refers to in presence of categorical. covariates).
  - `logLik` to output the log-likelihood of the estimated model.
  - `model.tables` to extract the estimates, standard errors, p-value, and confidence intervals.
  - `plot` to obtain a diagnostic plots, partial residual plots, or a graphical display of the fitted values.
  - `predict` to compute the mean conditional on covariates and possible outcome values.
  - `profile` to display the likelihood or profile likelihood of the model.
  - `resample` to use non-parametric bootstrap or permutation test for statistical inference.
  - `residuals` to extract the observed residuals of the fitted model, possibly normalized ($\widehat{\Omega}^{-\frac{1}{2}}\widehat{\varepsilon}$).
  - `sigma` to extract the modeled residual variance covariance matrix ($\widehat{\Omega}$).
  - `summary` to obtain a summary of the input, model fit, and estimated values.
  - `vcov` to extract the variance-covariance matrix of the mean parameters ($\widehat{\Sigma}_{\widehat{\beta}}$).

- the `mlmm` function to fit group-specific linear mixed models and gather the estimated coefficients.

- the `sampleRem` function to simulate longitudinal data.

- the `LMMstar.options` function enables the user to display the default values used in the **LMMstar** package. The function can also change the default values to better match the user needs.

1 Illustrative dataset

To illustrate the functionalities of the package, we will use the gastricbypass dataset. The long format can be imported using:

```
data(gastricbypassL, package = "LMMstar")
head(gastricbypassL)
```

```
  id visit time weight glucagonAUC
1 1     1  -13  127.2      20.690
2 2     1  -13  165.2      49.922
3 3     1  -13  109.7      42.434
4 4     1  -13  146.2      27.517
5 5     1  -13  113.1      29.151
6 6     1  -13  158.8      42.700
```

See `?gastricbypassL` for a presentation of the dataset. It is convenient to encode the time variable in two formats:

- numeric, e.g. here with the time in week since surgery (`time` variable taking values -13,-1,1,13 - negative times refering to before an intervention and positive times after the intervention).

- factor, e.g. here with the visit index (`visit` variable taking value 1,2,3,4)

To illustrate certain functionalities we will use an (artificial) group variable:

```
gastricbypassL$group <- as.factor(as.numeric(gastricbypassL$id)%%2)
```

and dichotomize time as before and after the intervention:

```
gastricbypassL$baseline <- gastricbypassL$time<0
```

The corresponding wide format is

```
data(gastricbypassW, package = "LMMstar")
head(gastricbypassW)
```

```
  id weight1 weight2 weight3 weight4 glucagonAUC1 glucagonAUC2 glucagonAUC3 glucagonAUC4
1 1   127.2   120.7   115.5   108.1       20.690       20.535       92.600       43.434
2 2   165.2   153.4   149.2   132.0       49.922       58.513       49.633       35.747
3 3   109.7   101.6    97.7    87.1       42.434       25.770       91.240       83.137
4 4   146.2   142.4   136.7   123.0       27.517       27.552       59.360       21.371
5 5   113.1   105.6    99.9    87.7       29.151           NA       86.859       57.970
6 6   158.8   143.6   134.6   108.7       42.700       31.616       53.408       37.636
```

for which we can also add the group variable:

```
gastricbypassW$group <- as.numeric(gastricbypassW$id)%%2
```

In some cases we will, for comparison, perform complete case analyses with the following dataset:

```
gastricbypassL.NNA <- gastricbypassL[!is.na(gastricbypassL$glucagonAUC),]
```

2 Visualization & descriptive statistics

2.1 Graphical display

A scatterplot of the data can obtained by specifying which columns to display when using the wide format:

```
scatterplot(gastricbypassW, ## left panel
            columns = c("weight1","weight2","weight3","weight4"))
```

When using the long format, a formula should describe the structure of the data: `outcome ~ order|cluster`

- the left hand side indicates the values to be displayed (here weight)

- the right hand side indicates the ordering of the repetitions (here over time) and how the repetitions are grouped within clusters (here within subject).

When calling `scatterplot`, the argument `group` leads to different color per group and the argument `type.diag` enables to use histograms (or density plots) instead of boxplots:

```
scatterplot(weight~time|id, data = gastricbypassL, ## right panel
            type.diag = "hist", group = "group")
```


By default the resulting object will be of class `list`. A `ggplot2` object can be obtained by setting the argument `facet` to `"grid2"`. This requires to have installed the package ggh4x and will produce a slightly different graphical display.

There is (currently) not dedicated function to obtain spaghetti plots. Instead one can use the ggplot2 package with the long format, e.g.:

```
gg.spa <- ggplot(gastricbypassL, aes(x=time,y=weight,group=id,color=id))
gg.spa <- gg.spa + geom_point() + geom_line()
gg.spa
```

2.2 Missing data patterns

The `summarizeNA` function identifies the possible combinations of observed/missing data:

```
mp <- summarizeNA(gastricbypassL)
mp
```

frequency	missing.pattern	n.missing	id	visit	time	weight	glucagonAUC	group	baseline
78	0000000	0	0	0	0	0	0	0	0
2	0000100	1	0	0	0	0	1	0	0

A graphical representation can be obtained using `plot`:

```
plot(mp)
```

See `help(plot.summarizeNA)` for options to customize the graphical display.

2.3 Summary statistics

Mean, standard deviation, and other summary statistic can be computed with respect to a categorical variable (typically time) using the `summarize` function:
(⚠ this function has the same name as a function from the dplyr package. If you have loaded dplyr, you should use `LMMstar:::summarize`)

```
sss <- summarize(weight+glucagonAUC ~ time, data = gastricbypassL, na.rm = TRUE)
print(sss, digits = 3)
```

	outcome	time	observed	missing	mean	sd	min	q1	median	q3	max
1	weight	-13	20	0	129.0	20.3	100.90	115.3	123.1	139.8	173.0
2		-1	20	0	121.2	18.9	95.70	107.8	114.5	134.5	162.2
3		1	20	0	115.7	18.3	89.90	102.2	110.6	128.4	155.0
4		13	20	0	102.4	17.1	78.80	90.4	98.5	108.2	148.0
5	glucagonAUC	-13	20	0	32.3	15.5	10.28	21.3	27.9	42.5	69.1
6		-1	19	1	29.7	13.7	9.87	21.2	25.8	33.6	67.7
7		1	19	1	76.9	27.9	35.85	56.5	73.8	91.9	135.9
8		13	20	0	52.0	21.0	21.37	37.2	51.2	57.9	109.2

Specifying a cluster (`id`) and ordering variable (`time`) enable to output correlation matrices:
(⚠ there should be no duplicated value of the ordering variable within cluster)

```
sss2 <- summarize(weight ~ time|id, data = gastricbypassL, na.rm = TRUE)
print(sss2, digits = 3)
```

|	time	observed	missing	mean	sd	min	q1	median	q3	max
1	-13	20	0	129	20.3	100.9	115.3	123.1	140	173
2	-1	20	0	121	18.9	95.7	107.8	114.5	135	162
3	1	20	0	116	18.3	89.9	102.2	110.6	128	155
4	13	20	0	102	17.1	78.8	90.4	98.5	108	148

```
 Pearson's correlation:
      -13    -1     1    13
-13 1.000 0.990 0.986 0.946
-1  0.990 1.000 0.997 0.959
1   0.986 0.997 1.000 0.966
13  0.946 0.959 0.966 1.000
```

Graphical displays of the summary statistics can be obtained via the `plot` method, where the argument `type` specifies the summary statistic to be displayed:

```
plot(sss2, type = "mean") ## left panel
plot(sss2, type = "sd") ## middle panel
plot(sss2, type = "cor") ## right panel
```

See `help(plot.summarize)` for options to customize the graphical display.

2.4 Correlation and partial correlations

The `partialCor` function can be used to evaluate group-specific correlations, e.g.:

```
partialCor(weight + glucagonAUC ~ 1, by = "group", data = gastricbypassL)
```

```
                          estimate    se   df  lower     upper p.value
0: rho(weight,glucagonAUC)  -0.328 0.143 21.8 -0.587 -0.00886  0.0447
1: rho(weight,glucagonAUC)  -0.354 0.141 22.5 -0.607 -0.03631  0.0313
```

This willl lead to the same estimate as the `cor.test` function (Pearson correlation):

```
gastricbypassL.0 <- gastricbypassL[gastricbypassL$group==0,]
rho <- cor.test(gastricbypassL.0$weight, gastricbypassL.0$glucagonAUC)
c(rho$estimate, p.value = rho$p.value)
```

```
      cor    p.value
-0.328481   0.038505
```

However the p-value may differ, especially in small samples, as `partialCor` uses a different (and probably more crude) small sample approximation for the estimator's distribution. Nevertheless `partialCor` enables to compare correlation coefficients across groups, by specifying the argument `effects`:

```
partialCor(weight + glucagonAUC ~ 1, by = "group", effects = "Dunnett",
           data = gastricbypassL)
```

```
                                                        estimate se df lower upper p.value
1:rho(weight,glucagonAUC) - 0:rho(weight,glucagonAUC)    -0.0255 NA NA    NA    NA   0.899
```

Partial correlations can be also computed by specifying covariate to adjust for on the right-hand side:

```
partialCor(weight4 + glucagonAUC4 ~ weight1,
           data = gastricbypassW)
```

```
                        estimate    se   df  lower upper p.value
rho(weight4,glucagonAUC4)   0.112 0.233 9.12 -0.397 0.568   0.645
```

When the set of covariates is outcome-dependent, a list of formulas can be used instead:

```
partialCor(list(weight1 ~ glucagonAUC1, weight4 ~ glucagonAUC4),
           data = gastricbypassW)
```

```
                    estimate     se   df lower upper  p.value
rho(weight1,weight4)   0.946 0.0252 26.4 0.861 0.979 5.51e-08
```

These partial correlations are defined as the residual correlation between the outcomes, i.e. the correlation once the covariate effects have been substracted from the outcome, and a linear mixed model is used to estimated them.

3 Multiple Student's t-tests

When working with multiple outcomes and having no missing data, mean comparisons between exposure groups can be carried out using Student's t-tests at each timepoint, e.g.:

```
restt <- t.test(weight1 ~ group, data = gastricbypassW)
c(estimate = unname(diff(restt$estimate)), p.value = restt$p.value)
```

```
 estimate    p.value
-10.60000    0.25282
```

And so on for the three other timepoints. Morever results would typically need to be adjusted for multiple comparisons, e.g. when looking for any mean difference. This can be faciliated by

```
## single step max-test adjustment (see help(confint.Wald_lmm) for details)
mt.test(weight1+weight2+weight3+weight4~group, data = gastricbypassW)
```

```
       by parameter estimate     se     df   lower   upper p.value
1 weight1     group   -10.60 8.9717 17.965 -30.968  9.7680 0.31894
2 weight2     group    -9.50 8.3951 17.985 -28.559  9.5590 0.34164
3 weight3     group    -8.92 8.1295 17.959 -27.376  9.5358 0.35891
4 weight4     group    -4.59 7.7607 17.682 -22.209 13.0286 0.66331
```

The method used to adjust confidence intervals and p-values for multiple comparisons can be specified via the `method` argument, e.g.:

```
## no adjustment
mt.test(weight1+weight2+weight3+weight4~group, data = gastricbypassW, method = "none")
```

```
       by parameter estimate     se     df   lower   upper p.value
1 weight1     group   -10.60 8.9717 17.965 -29.452  8.2516 0.25281
2 weight2     group    -9.50 8.3951 17.985 -27.139  8.1386 0.27266
3 weight3     group    -8.92 8.1295 17.959 -26.002  8.1622 0.28703
4 weight4     group    -4.59 7.7607 17.682 -20.916 11.7356 0.56171
```

```
## bonferroni adjustment
mt.test(weight1+weight2+weight3+weight4~group, data = gastricbypassW, method = "bonferroni")
```

```
       by parameter estimate     se     df   lower  upper p.value
1 weight1     group   -10.60 8.9717 17.965 -35.498 14.298       1
2 weight2     group    -9.50 8.3951 17.985 -32.795 13.795       1
3 weight3     group    -8.92 8.1295 17.959 -31.481 13.641       1
4 weight4     group    -4.59 7.7607 17.682 -26.165 16.985       1
```

4 Linear mixed model (LMM)

4.1 Classical covariance patterns

Several build-in covariance patterns can be used when specifying the linear model. The most basic ones are the **identity** structure:

```
eId.lmm <- lmm(glucagonAUC ~ visit*group, repetition = ~time|id,
               structure = "ID", data = gastricbypassL)
eId.lmm
cat(" modeled residual variance-covariance: \n");sigma(eId.lmm)
```

```
                  Linear regression

 outcome/cluster/time: glucagonAUC/id/time
 data                 : 78 observations from 20 clusters
 parameter            : 8 mean ((Intercept) visit2 visit3 visit4 group1 visit2:group1 visit3:group1
                        1 variance (sigma)
 log-restr.likelihood: -316.461119970244
 convergence          : TRUE (0 iterations)
 modeled residual variance-covariance:
         -13      -1       1       13
-13  381.35    0.00    0.00    0.00
-1     0.00  381.35    0.00    0.00
1      0.00    0.00  381.35    0.00
13     0.00    0.00    0.00  381.35
```

and the **independence** structure:

```
eInd.lmm <- lmm(glucagonAUC ~ visit*group, repetition = ~time|id,
               structure = "IND", data = gastricbypassL)
eInd.lmm
cat(" modeled residual variance-covariance: \n");sigma(eInd.lmm)
```

```
           Linear regression with heterogeneous residual variance

 outcome/cluster/time: glucagonAUC/id/time
 data                 : 78 observations from 20 clusters
 parameter            : 8 mean ((Intercept) visit2 visit3 visit4 group1 visit2:group1 visit3:group1
                        4 variance (sigma k.-1 k.1 k.13)
 log-restr.likelihood: -310.428096419287
 convergence          : TRUE (0 iterations)
 modeled residual variance-covariance:
         -13      -1       1       13
-13  209.44    0.00    0.00    0.00
-1     0.00  174.81    0.00    0.00
1      0.00    0.00  768.23    0.00
13     0.00    0.00    0.00  382.95
```

The most common linear mixed model uses a **compound symmetry** structure:

```
eCS.lmm <- lmm(glucagonAUC ~ visit*group, repetition = ~time|id,
               structure = "CS", data = gastricbypassL)
eCS.lmm
cat(" modeled residual variance-covariance: \n");sigma(eCS.lmm)
```

```
            Linear Mixed Model with a compound symmetry covariance matrix

 outcome/cluster/time: glucagonAUC/id/time
 data                 : 78 observations from 20 clusters
 parameter            : 8 mean ((Intercept) visit2 visit3 visit4 group1 visit2:group1 visit3:group1
                        1 variance (sigma)
                        1 correlation (rho(id))
 log-restr.likelihood: -314.394203759159
 convergence          : TRUE (6 iterations)
 modeled residual variance-covariance:
        -13      -1       1      13
-13 380.580  82.741  82.741  82.741
-1   82.741 380.580  82.741  82.741
1    82.741  82.741 380.580  82.741
13   82.741  82.741  82.741 380.580
```

A more flexible model can be obtained with a **toeplitz** covariance matrix:

```
eTOE.lmm <- lmm(glucagonAUC ~ visit*group, repetition = ~time|id,
               structure = "TOEPLITZ", data = gastricbypassL)
eTOE.lmm
cat(" modeled residual correlation: \n");cov2cor(sigma(eTOE.lmm))
```

```
            Linear Mixed Model with a block Toeplitz covariance matrix

 outcome/cluster/time: glucagonAUC/id/time
 data                 : 78 observations from 20 clusters
 parameter            : 8 mean ((Intercept) visit2 visit3 visit4 group1 visit2:group1 visit3:group1
                        4 variance (sigma k.-1 k.1 k.13)
                        4 correlation (rho(12) rho(14) rho(26) rho(2))
 log-restr.likelihood: -297.525485582536
 convergence          : TRUE (15 iterations)
 modeled residual correlation:
          -13       -1        1        13
-13  1.000000 0.700020 0.093615 -0.082963
-1   0.700020 1.000000 0.016795  0.093615
1    0.093615 0.016795 1.000000  0.700020
13  -0.082963 0.093615 0.700020  1.000000
```

And an even more flexible model can be obtained with an **unstructured** covariance matrix:

```
eUN.lmm <- lmm(glucagonAUC ~ visit*group, repetition = ~time|id,
               structure = "UN", data = gastricbypassL)
eUN.lmm
cat(" modeled residual variance-covariance: \n");sigma(eUN.lmm)
```

```
                Linear Mixed Model with an unstructured covariance matrix

 outcome/cluster/time: glucagonAUC/id/time
 data                 : 78 observations from 20 clusters
 parameter            : 8 mean ((Intercept) visit2 visit3 visit4 group1 visit2:group1 visit3:group1
                        4 variance (sigma k.-1 k.1 k.13)
                        6 correlation (rho(-13,-1) rho(-13,1) rho(-13,13) rho(-1,1) rho(-1,13) rho(1
 log-restr.likelihood: -295.314056198772
 convergence          : TRUE (8 iterations)
 modeled residual variance-covariance:
         -13        -1         1        13
-13 209.442 150.2502 106.4000 -24.202
-1  150.250 168.1138   1.3064 -23.884
1   106.400   1.3064 748.0769 288.184
13  -24.202 -23.8844 288.1839 382.952
```

Stratification of the covariance structure on a categorical variable is also possible:

- e.g. to get a **stratified compound symmetry**

```
eSCS.lmm <- lmm(glucagonAUC ~ visit*group, repetition = ~time|id,
                structure = CS(group~1), data = gastricbypassL)
eSCS.lmm
```

```
              Linear Mixed Model with a stratified compound symmetry covariance matrix

outcome/cluster/time: glucagonAUC/id/time
data                 : 78 observations from 20 clusters
parameter            : 8 mean ((Intercept) visit2 visit3 visit4 group1 visit2:group1 visit3:group1 v
                        2 variance (sigma:0 sigma:1)
                        2 correlation (rho(id):0 rho(id):1)
log-restr.likelihood: -314.123797063042
convergence          : TRUE (7 iterations)
```

- e.g. **stratified unstructured** covariance matrix:

```
eSUN.lmm <- lmm(glucagonAUC ~ visit*group, repetition = ~time|id,
                structure = UN(group~1), data = gastricbypassL)
eSUN.lmm
```

Linear Mixed Model with a stratified unstructured covariance matrix

```
outcome/cluster/time: glucagonAUC/id/time
data                : 78 observations from 20 clusters
parameter           : 8 mean ((Intercept) visit2 visit3 visit4 group1 visit2:group1 visit3:group1 v
                      8 variance (sigma:0 sigma:1 k.-1:0 k.1:0 k.13:0 k.-1:1 k.1:1 k.13:1)
                      12 correlation (rho(-13,-1):0 rho(-13,1):0 rho(-13,13):0 rho(-1,1):0 rho(-1,1
log-restr.likelihood: -286.536815485471
convergence         : TRUE (10 iterations)
```

with modeled residual variance-covariance:

sigma(eSCS.lmm)	sigma(eSUN.lmm)

```
 $`0`                                      $`0`
          -13      -1       1      13                -13      -1       1      13
 -13  334.289  50.782  50.782  50.782       -13  309.85 251.512 102.189 -42.250
 -1    50.782 334.289  50.782  50.782       -1   251.51 274.752 -79.811 -90.718
 1     50.782  50.782 334.289  50.782       1    102.19 -79.811 579.110 163.767
 13    50.782  50.782  50.782 334.289       13   -42.25 -90.718 163.767 173.439


 $`1`                                      $`1`
          -13      -1       1      13                 -13      -1       1       13
 -13  428.46  115.09  115.09  115.09       -13  109.0309 48.667 104.908  -6.1549
 -1   115.09  428.46  115.09  115.09       -1    48.6665 59.395  93.976  43.2144
 1    115.09  115.09  428.46  115.09       1    104.9077 93.976 967.583 450.8899
 13   115.09  115.09  115.09  428.46       13    -6.1549 43.214 450.890 592.4655
```

Finally the some covariance patterns like the compound symmetry structure may depend on covariates:

- e.g. to obtain a **block compound symmetry** structure[1]:

```
eBCS.lmm <- lmm(glucagonAUC ~ visit*group, repetition = ~time|id,
                structure = CS(~baseline, type = "homogeneous"), data = gastricbypassL)
eBCS.lmm
cat(" modeled residual variance-covariance: \n");sigma(eBCS.lmm)
```

```
                  Linear Mixed Model with a block compound symmetry covariance matrix

 outcome/cluster/time: glucagonAUC/id/time
 data               : 78 observations from 20 clusters
 parameter          : 8 mean ((Intercept) visit2 visit3 visit4 group1 visit2:group1 visit3:group1
                      1 variance (sigma)
                      2 correlation (rho(id/baseline) rho(id))
 log-restr.likelihood: -308.994835006264
 convergence         : TRUE (6 iterations)
 modeled residual variance-covariance:
         -13      -1       1       13
-13 380.957 226.403  15.465  15.465
-1  226.403 380.957  15.465  15.465
1    15.465  15.465 380.957 226.403
13   15.465  15.465 226.403 380.957
```

- e.g. to obtain a **block unstructured** covariance matrix:

```
eBUN.lmm <- lmm(glucagonAUC ~ visit*group, repetition = ~time|id,
                structure = CS(~baseline, type = "heterogeneous"), data = gastricbypassL)
eBUN.lmm
cat(" modeled residual variance-covariance: \n");sigma(eBUN.lmm)
```

```
                  Linear Mixed Model with a block unstructured covariance matrix

 outcome/cluster/time: glucagonAUC/id/time
 data                 : 78 observations from 20 clusters
 parameter            : 8 mean ((Intercept) visit2 visit3 visit4 group1 visit2:group1 visit3:group1
                        2 variance (sigma k.TRUE)
                        3 correlation (rho(FALSE) rho(FALSE,TRUE) rho(TRUE))
 log-restr.likelihood: -300.047474124556
 convergence         : TRUE (7 iterations)
 modeled residual variance-covariance:
         -13      -1       1       13
-13 189.420 150.356  15.353  15.353
-1  150.356 189.420  15.353  15.353
1    15.353  15.353 570.908 300.071
13   15.353  15.353 300.071 570.908
```

---

[1]similar to nested random effects

4.2 User-specific covariance patterns

It is possible input user-specific covariance patterns under the following model for the residuals:

$$\Omega = \boldsymbol{\sigma}^{\mathsf{T}} R \boldsymbol{\sigma}$$

- $\boldsymbol{\sigma} = f(\boldsymbol{\theta}_\sigma, Z_\sigma)$ is a vector of residual standard errors depending on a vector of parameters $\boldsymbol{\theta}_\sigma$ and possible covariates via the design matrix Z_σ.

- $R = g(\boldsymbol{\theta}_R, Z_R)$ is a matrix of residual correlations depending on a vector of parameters $\boldsymbol{\theta}_R$ and possible covariates via the design matrix Z_R.

To be more concrete, consider the following correlation matrix

```
rho.2block <- function(p,n.time,X){
  rho <- matrix(1, nrow = n.time, ncol = n.time)
  rho[1,2] <- rho[2,1] <- rho[4,5] <- rho[5,4] <- p["rho1"]
  rho[1,3] <- rho[3,1] <- rho[4,6] <- rho[6,4] <- p["rho2"]
  rho[2,3] <- rho[3,2] <- rho[5,6] <- rho[6,5] <- p["rho3"]
  rho[4:6,1:3] <- rho[1:3,4:6] <- p["rho4"]
  return(rho)
}
Rho <- rho.2block(p = c(rho1=0.25,rho2=0.5,rho3=0.4,rho4=0.1),
                  n.time = 6)
Rho
```

```
      [,1] [,2] [,3] [,4] [,5] [,6]
[1,] 1.00 0.25  0.5 0.10 0.10  0.1
[2,] 0.25 1.00  0.4 0.10 0.10  0.1
[3,] 0.50 0.40  1.0 0.10 0.10  0.1
[4,] 0.10 0.10  0.1 1.00 0.25  0.5
[5,] 0.10 0.10  0.1 0.25 1.00  0.4
[6,] 0.10 0.10  0.1 0.50 0.40  1.0
```

and the corresponding dataset:

```
set.seed(11)
Y <- mvtnorm::rmvnorm(1000, mean = rep(0,6), sigma = Rho)
dfW <- cbind(id = 1:NROW(Y), as.data.frame(Y))
dfL <- reshape2::melt(dfW, id.vars = "id", variable.name = "time")
dfL[dfL$id %in% 1:2,]
```

```
        id time      value                      id time      value
   1     1   V1 -0.9842079            2     2   V1  1.2402726
   1001  1   V2 -0.3681245            1002  2   V2  0.6494215
   2001  1   V3 -1.6174652            2002  2   V3  0.3272105
   3001  1   V4 -1.4994103            3002  2   V4 -1.0626973
   4001  1   V5  0.7493107            4002  2   V5 -0.9013244
   5001  1   V6 -1.0719657            5002  2   V6 -0.6696714
```

To estimate the corresponding mixed model we first define a new covariance structure:

```
myStruct <- CUSTOM(~time,
                   FCT.sigma = function(p,n.time,X){rep(p,n.time)}, ## function f
                   init.sigma = c("sigma"=1),
                   FCT.rho = rho.2block, ## function g
                   init.rho = c("rho1"=0.25,"rho2"=0.25,"rho3"=0.25,"rho4"=0.25))
```

and then call `lmm` with this structure structure:

```
e.lmmCUSTOM <- lmm(value~time, repetition=~time|id,
                   structure = myStruct, data=dfL,
                   df = FALSE) ## df = FALSE to save computation time
logLik(e.lmmCUSTOM)
```

```
[1] -7962.243
```

The optimization procedure may be slow but should eventually reaches an optimum. We can then output the estimated correlation matrix:

```
cov2cor(sigma(e.lmmCUSTOM))
```

```
           V1         V2         V3         V4         V5         V6
V1 1.00000000 0.24898095 0.50058994 0.09053785 0.09053785 0.09053785
V2 0.24898095 1.00000000 0.36110943 0.09053785 0.09053785 0.09053785
V3 0.50058994 0.36110943 1.00000000 0.09053785 0.09053785 0.09053785
V4 0.09053785 0.09053785 0.09053785 1.00000000 0.24898095 0.50058994
V5 0.09053785 0.09053785 0.09053785 0.24898095 1.00000000 0.36110943
V6 0.09053785 0.09053785 0.09053785 0.50058994 0.36110943 1.00000000
```

Comparison to build-in structure: consider the following model using a build-in compound symmetry structure:

```
system.time(
  e.lmmDEFAULT.CS <- lmm(value~time, repetition = ~time|id,
                         structure = "CS", data = dfL,
                         df = FALSE)
)
```

```
 user   system elapsed
0.097    0.000   0.097
```

Using instead `CUSTOM` to specifying this structure:

```
myCS <- CUSTOM(~1,
               FCT.sigma = function(p,n.time,X){rep(p,n.time)},
               init.sigma = c("sigma"=1),
               FCT.rho = function(p,n.time,X){p+diag(1-p,n.time,n.time)},
               init.rho = c("rho"=0.5))
```

16

is considerably slower than using the pre-specified structure:

```
system.time(
  e.lmmCUSTOM.CS <- lmm(value~time, repetition = ~time|id,
                        structure = myCS, data = dfL,
                        df = FALSE)
)
```

```
 user   system  elapsed
0.952    0.019    0.972
```

but will lead to the same estimates:

```
logLik(e.lmmDEFAULT.CS)
logLik(e.lmmCUSTOM.CS)
```

```
[1] -8186.859
[1] -8186.859
```

There are two reasons for the slower execution time: slower evaluation of the derivatives (since they are obtained by numerical differentiation) and worse starting point, as reflected by the larger number of interations needed to reach convergence:

```
e.lmmDEFAULT.CS$opt$n.iter
e.lmmCUSTOM.CS$opt$n.iter
```

```
[1] 1
[1] 4
```

Faster execution time can be obtained by specifying the first and second derivative regarding each parameter:

```
myCS.wD <- CUSTOM(~1,
                  FCT.sigma = function(p,n.time,X){rep(p,n.time)},
                  dFCT.sigma = function(p,n.time,X){list(sigma = rep(1,n.time))},
                  d2FCT.sigma = function(p,n.time,X){list(sigma = rep(0,n.time))},
                  init.sigma = c("sigma"=1),
                  FCT.rho = function(p,n.time,X){p+diag(1-p,n.time,n.time)},
                  dFCT.rho = function(p,n.time,X){list(rho = 1-diag(1,n.time,n.time))},
                  d2FCT.rho = function(p,n.time,X){list(rho = matrix(0,n.time,n.time))},
                  init.rho = c("rho"=0.5))
```

```
system.time(
  e.lmmCUSTOMwD.CS <- lmm(value~time,
                          repetition = ~time|id,
                          structure = myCS.wD,
                          data = dfL, df = FALSE
                          )
)
```

```
 user   system  elapsed
0.699    0.004    0.703
```

4.3 Estimation procedure

Initialiation: by default the mean parameters are initialized using Ordinary Least Squares (OLS) and the variance and correlation parameters are initialized by minimizing the difference between the observed and residuals variance-covariance matrix. These values can be visualized by specifying the argument `control`:

```
eCS.lmm.bis <- update(eCS.lmm, control = list(trace = 2))
```

```
Initialization:
  (Intercept)         visit2         visit3         visit4         group1 visit2:group1 visit3:group1
     38.72897       -4.73433       31.43303        4.52138      -12.82462        3.75946       27.00150
visit4:group1          sigma        rho(id)
     30.22391       19.52828        0.22819

Loop:
******
  (Intercept)         visit2         visit3         visit4         group1 visit2:group1 visit3:group1
     38.72897       -4.73433       31.43303        4.52138      -12.82462        3.80337       27.48103
visit4:group1          sigma        rho(id)
     30.22391       19.50846        0.21741
Convergence after 6 iterations: max score=1.2413e-05 | max change in coefficient=4.5167e-06
```

It is possible to input user-defined value:

- for all parameters (vector)

```
init.all <- coef(eCS.lmm, effects = "all")
eCS.lmm.bis <- update(eCS.lmm, control = list(init = init.all, trace = 1))
```

```
Convergence after 0 iteration: max score=1.2413e-05
```

- the mean parameters only (vector)

```
init.mean <- coef(eCS.lmm, effects = "mean")
eCS.lmm.bis <- update(eCS.lmm, control = list(init = init.mean, trace = 2))
```

```
Initialization:
  (Intercept)         visit2         visit3         visit4         group1 visit2:group1 visit3:group1
     38.72897       -4.73433       31.43303        4.52138      -12.82462        3.80337       27.48103
visit4:group1          sigma        rho(id)
     30.22391       19.52904        0.22849

Loop:
******
  (Intercept)         visit2         visit3         visit4         group1 visit2:group1 visit3:group1
     38.72897       -4.73433       31.43303        4.52138      -12.82462        3.80337       27.48103
visit4:group1          sigma        rho(id)
     30.22391       19.50846        0.21741
Convergence after 6 iterations: max score=1.4893e-05 | max change in coefficient=5.3866e-06
```

- a full data variance-covariance matrix (matrix).

```
init.vcov <- sigma(eCS.lmm)
eCS.lmm.bis <- update(eCS.lmm, control = list(init = init.vcov, trace = 1))
```

Convergence after 0 iteration: max score=1.2413e-05

Optimizer: by default the optimizer is a Newton Raphson algorithm with backtracking. At each iteration:

- it computes the first two moments (score, information) according to the current parameters values.

- it updates the variance-covariance parameters according to the gradient multiplied by the inverse of the information.

- it updates the mean parameters by generalized least squares (using the updated variance-covariance parameters).

- it checks whether the log-likelihoood at the updated estimates is well defined and higher than at the previous estimates. If this is not the case, the step is re-run with half the update of the variance-covariance parameters (backtracking).

One can modify the maximum number of iterations (`n.iter`), maximum number of backtracking steps (`n.backtracking`), the maximum score (absolute) value over all parameters (`tol.score`) and (absolute) maximum difference in parameter value between to iterations (`tol.param`) used to declare convergence. It is also possible to use another optimizer (`optimizer`). All these elements should be passed to the argument `control` of `lmm` using a list.

4.4 Model output

The summary method can be used to display the main information relative to the model fit:

```
summary(eUN.lmm)
```

```
              Linear Mixed Model

Dataset: gastricbypassL

  - 20 clusters
  - 78 observations were analyzed, 2 were excluded because of missing values
  - between 3 and 4 observations per cluster

Summary of the outcome and covariates:

    $ glucagonAUC: num  20.7 49.9 42.4 27.5 29.2 ...
    $ visit      : Factor w/ 4 levels "1","2","3","4": 1 1 1 1 1 1 1 1 1 1 ...
    $ group      : Factor w/ 2 levels "0","1": 2 1 2 1 2 1 2 1 2 1 ...
    reference level: visit=1;group=0

Estimation procedure

  - Restricted Maximum Likelihood (REML)
  - log-likelihood :-295.31
  - parameters: mean = 8, variance = 4, correlation = 6
  - convergence: TRUE (8 iterations)
    largest |score| = 4.6771e-05 for rho(-1,1)
            |change|= 1.68033723859651e-05 for visit3:group1

Residual variance-covariance: unstructured

  - correlation structure: ~0 + time
          -13       -1       1       13
  -13  1.0000   0.80072 0.26880 -0.0855
  -1   0.8007   1.00000 0.00368 -0.0941
  1    0.2688   0.00368 1.00000  0.5384
  13  -0.0855  -0.09413 0.53842  1.0000

  - variance structure: ~time
            standard.deviation ratio
  sigma.-13               14.5 1.000
  sigma.-1                13.0 0.896
  sigma.1                 27.4 1.890
  sigma.13                19.6 1.352
```

```
Fixed effects: glucagonAUC ~ visit * group

                 estimate     se     df   lower  upper p.value
   (Intercept)     38.729  4.576     18  29.114 48.344 < 1e-04 ***
   visit2          -4.734  2.776   17.5 -10.577  1.109 0.10574
   visit3          31.433   8.63   17.6  13.272 49.594 0.00192   **
   visit4           4.521  8.005     18 -12.297  21.34 0.57917
   group1         -12.825  6.472     18 -26.422  0.773 0.06302    .
   visit2:group1    3.987  3.996   17.9   -4.41 12.383 0.33169
   visit3:group1   27.571  12.42   17.8   1.461 53.682 0.03963    *
   visit4:group1   30.224 11.321     18   6.439 54.008 0.01562    *
   ------------------------------------------------------
   Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1.
   Columns lower and upper contain 95% pointwise confidence intervals for each coefficient.
   Model-based standard errors are derived from the observed information (column se).
   Degrees of freedom were computed using a Satterthwaite approximation (column df).
```

Note: the calculation of the degrees of freedom, especially when using the observed information can be quite slow. Setting the arguments `df` to `FALSE` and `type.information` to `"expected"` when calling `lmm` should lead to a more reasonnable computation time.

4.5   Extract estimated coefficients

The value of the estimated coefficients can be output using `coef`:

```
coef(eUN.lmm)
```

```
  (Intercept)         visit2         visit3         visit4         group1 visit2:group1 visit3:group1
      38.7290        -4.7343        31.4330         4.5214       -12.8246        3.9866       27.5714
visit4:group1
      30.2239
```

Variance coefficients can be output by specifying the `effects` argument:

```
coef(eUN.lmm, effects = "variance")
```

```
   sigma      k.-1       k.1       k.13
14.47212   0.89592   1.88991   1.35220
```

The first coefficient is the residual standard deviation at the reference timepoint (here -13 week) and the remaining coefficient the residual standard deviation at later timepoints relative to the reference timepoint. It is possible to apply specific transformation on the variance coefficients, for instance to obtain the residual variance at each timepoint:

```
coef(eUN.lmm, effects = "variance", transform.k = "sd")
```

```
sigma.-13  sigma.-1   sigma.1   sigma.13
   14.472    12.966    27.351    19.569
```

4.6   Extract estimated coefficient and associated uncertainty

The uncertainty about the mean coefficients can be obtained using the `model.tables` method [2]:

```
model.tables(eUN.lmm)
```

	estimate	se	df	lower	upper	p.value
(Intercept)	38.7290	4.5765	18.003	29.1143	48.34369	1.0891e-07
visit2	-4.7343	2.7759	17.543	-10.5772	1.10851	1.0574e-01
visit3	31.4330	8.6297	17.585	13.2719	49.59411	1.9229e-03
visit4	4.5214	8.0050	17.995	-12.2968	21.33958	5.7917e-01
group1	-12.8246	6.4721	18.003	-26.4219	0.77265	6.3015e-02
visit2:group1	3.9866	3.9957	17.937	-4.4102	12.38329	3.3169e-01
visit3:group1	27.5714	12.4199	17.831	1.4605	53.68232	3.9634e-02
visit4:group1	30.2239	11.3208	17.995	6.4394	54.00840	1.5624e-02

Values for the all correlation parameters can be displayed too, by specifying `effect=c("variance","correla`

```
model.tables(eUN.lmm, effect = c("variance","correlation"))
```

	estimate	se	df	lower	upper	p.value
sigma	14.4721183	2.412020	15.3158	10.15148	20.63170	NA
k.-1	0.8959206	0.127032	20.2671	0.66670	1.20396	0.44721963
k.1	1.8899095	0.431098	25.9157	1.18244	3.02067	0.00974152
k.13	1.3521979	0.317550	29.8074	0.83694	2.18468	0.20874407
rho(-13,-1)	0.8007214	0.085177	13.4142	0.52949	0.92343	0.00042923
rho(-13,1)	0.2688043	0.219200	7.9286	-0.26374	0.67576	0.27735748
rho(-13,13)	-0.0854578	0.233981	8.5882	-0.55306	0.42309	0.72505145
rho(-1,1)	0.0036838	0.237237	8.1487	-0.49424	0.49979	0.98798445
rho(-1,13)	-0.0941328	0.233649	8.9191	-0.55697	0.41331	0.69821381
rho(1,13)	0.5384239	0.176221	10.2233	0.05058	0.81883	0.03522642

Because these parameters are constrained (e.g. strictly positive), they uncertainty is by default computed after transformation (e.g. `log`) and then backtransformed. The column argument can be used to extract more or less information, e.g.:

```
model.tables(eUN.lmm, columns = c("estimate","p.value"))
```

	estimate	p.value
(Intercept)	38.7290	1.0891e-07
visit2	-4.7343	1.0574e-01
visit3	31.4330	1.9229e-03
visit4	4.5214	5.7917e-01
group1	-12.8246	6.3015e-02
visit2:group1	3.9866	3.3169e-01
visit3:group1	27.5714	3.9634e-02
visit4:group1	30.2239	1.5624e-02

---

[2]it is equivalent to `confint` method except that by default it also outputs `se` and `p.value`

All parameters can be displayed by specifying `effect="all"`. The functions `add` (resp. `remove`) can be used to add (resp. remove) one or several columns from the default display, e.g.:

```
model.tables(eUN.lmm, columns = add("statistic"))
```

```
              estimate       se  statistic      df     lower      upper    p.value
(Intercept)    38.7290   4.5765    8.46260  18.003   29.1143  48.34369  1.0891e-07
visit2         -4.7343   2.7759   -1.70552  17.543  -10.5772   1.10851  1.0574e-01
visit3         31.4330   8.6297    3.64242  17.585   13.2719  49.59411  1.9229e-03
visit4          4.5214   8.0050    0.56482  17.995  -12.2968  21.33958  5.7917e-01
group1        -12.8246   6.4721   -1.98151  18.003  -26.4219   0.77265  6.3015e-02
visit2:group1   3.9866   3.9957    0.99772  17.937   -4.4102  12.38329  3.3169e-01
visit3:group1  27.5714  12.4199    2.21995  17.831    1.4605  53.68232  3.9634e-02
visit4:group1  30.2239  11.3208    2.66977  17.995    6.4394  54.00840  1.5624e-02
```

4.7 Extract estimated residual variance-covariance structure

The method `sigma` can be used to output the modeled residual covariance structure and then converted to a correlation matrix using `cov2cor`:

```
Sigma <- sigma(eUN.lmm)
Sigma
```

```
cov2cor(Sigma)
```

```
          -13      -1       1       13
-13  209.442  150.250  106.400  -24.202
-1   150.250  168.114    1.306  -23.884
1    106.400    1.306  748.077  288.184
13   -24.202  -23.884  288.184  382.952
```

```
         -13      -1      1      13
-13    1.000   0.801  0.269  -0.085
-1     0.801   1.000  0.004  -0.094
1      0.269   0.004  1.000   0.538
13    -0.085  -0.094  0.538   1.000
```

The method can also be used to extract the residual covariance relative to a "known" individual:

```
sigma(eUN.lmm, cluster = 5)
```

```
        -13       1       13
-13  209.442  106.40  -24.202
1    106.400  748.08  288.184
13   -24.202  288.18  382.952
```

or for a new individual:

```
newdata <- data.frame(id = "X", time = c("-13","-1","1","13"))
sigma(eUN.lmm, cluster = newdata)
```

```
        -13       -1        1       13
-13  209.442  150.2502  106.4000  -24.202
-1   150.250  168.1138    1.3064  -23.884
1    106.400    1.3064  748.0769  288.184
13   -24.202  -23.8844  288.1839  382.952
```

4.8 Marginal effects

The `effects` method can be used to evaluate marginal means with respect to a categorical variable:

- $\mathbb{E}[Y_t \mid \text{group}]$

```
effects(eUN.lmm, variable = "group")
```

```
                Average counterfactual outcome
                w.r.t 'group' values

                   estimate    se    df  lower    upper
   group=0(t=-13)    38.729 4.576    18 29.114   48.344
   group=0(t=-1)     33.995   4.1  17.9 25.377   42.612
   group=0(t=1)      70.162 8.649  17.7 51.968   88.356
   group=0(t=13)      43.25 6.188    18 30.249   56.251
   group=1(t=-13)    25.904 4.576    18  16.29   35.519
   group=1(t=-1)     25.157 4.167  18.7 16.425   33.889
   group=1(t=1)      84.909 8.951  18.2 66.115  103.702
   group=1(t=13)      60.65 6.188    18 47.649   73.651
```

- $\mathbb{E}[Y_t - Y_0 \mid \text{group}]$

```
effects(eUN.lmm, type = "change", variable = "group")
```

```
                Average counterfactual change in outcome
                w.r.t 'group' values

                    estimate    se    df   lower   upper
   group=0(dt=-1)     -4.734 2.776  17.5 -10.577   1.109
   group=0(dt=1)      31.433  8.63  17.6  13.272  49.594
   group=0(dt=13)      4.521 8.005    18 -12.297   21.34
   group=1(dt=-1)     -0.748 2.874  18.3  -6.779   5.283
   group=1(dt=1)      59.004 8.932    18  40.242  77.767
   group=1(dt=13)     34.745 8.005    18  17.927  51.563
```

- $\mathbb{E}\left[\int_0^T Y_t dt \mid \text{group}\right]$

```
effects(eUN.lmm, type = "auc", variable = "group")
```

```
            Average counterfactual area under the outcome curve
             w.r.t 'group' values

              estimate      se   df     lower     upper
group=0(auc) 1220.972 104.098 17.8 1002.072  1439.873
group=1(auc) 1289.782 105.512 18.5 1068.508  1511.056
```

It can also be used to contrast these marginal means:

- $\mathbb{E}\left[Y_t \mid \text{group} = 1\right] - \mathbb{E}\left[Y_t \mid \text{group} = 0\right]$

```
effects(eUN.lmm, type = "difference", variable = "group")
```

```
            Difference in average counterfactual outcome
             w.r.t 'group' values


                 estimate     se   df   lower   upper p.value
group=1-0(t=-13)  -12.825  6.472   18 -26.422   0.773  0.0630 .
group=1-0(t=-1)    -8.838  5.846 18.3 -21.106    3.43  0.1477
group=1-0(t=1)     14.747 12.447 17.9 -11.409  40.903  0.2516
group=1-0(t=13)    17.399  8.752   18  -0.987  35.785  0.0622 .
```

- $\mathbb{E}\left[Y_t - Y_0 \mid \text{group} = 1\right] - \mathbb{E}\left[Y_t - Y_0 \mid \text{group} = 0\right]$

```
effects(eUN.lmm, type = c("change","difference"), variable = "group")
```

```
            Difference in average counterfactual change in outcome
             w.r.t 'group' values


                estimate     se   df lower  upper p.value
group=1-0(dt=-1)    3.987  3.996 17.9 -4.41 12.383  0.3317
group=1-0(dt=1)    27.571  12.42 17.8 1.461 53.682  0.0396 *
group=1-0(dt=13)   30.224 11.321   18 6.439 54.008  0.0156 *
```

- $\mathbb{E}\left[\int_0^T Y_t dt \mid \text{group} = 1\right] - \mathbb{E}\left[\int_0^T Y_t dt \mid \text{group} = 0\right]$

```
effects(eUN.lmm, type = c("auc","difference"), variable = "group")
```

```
            Difference in average counterfactual area under the outcome curve
             w.r.t 'group' values

               estimate     se   df   lower   upper p.value
group=1-0(auc)   68.809 148.22 18.1 -242.44 380.059   0.648
```

It is possible to control the set of covariates used to condition on via the `conditional` argument. This can be useful when considering an interaction with a biomarker to obtain biomarker-specific effects.

4.9 Random effects

Mixed model having a compound symmetry structure with positive correlation parameters may be equivalent to random intercept models, possibly with nested random effects. Indeed in some case the residual variance-covariance matrix can then be decomposed as:

$$\Omega = Z\Psi Z^{\mathsf{T}} + \Delta$$

- Z is the design matrix associated to the random effect (e.g. patient id)

- Ψ is the variance-covariance of the random effects

- Δ the residual variance covariance conditional to the random effects.

One can the use `lme4` syntax to fit random intercept models with `lmm`:

```
eRI.lmm <- lmm(glucagonAUC ~ visit*group + (1|id), data = gastricbypassL)
eRI.lmm
```

```
            Linear Mixed Model with a random intercept

outcome/cluster/time: glucagonAUC/id/XXtimeXX
data                 : 78 observations from 20 clusters
parameter            : 8 mean ((Intercept) visit2 visit3 visit4 group1 visit2:group1 visit3:group1 v
                       1 variance (sigma)
                       1 correlation (rho(id))
log-restr.likelihood: -314.394203759159
convergence          : TRUE (6 iterations)
```

It is also possible to specify cross or nested random effects, e.g.:

```
eNRI.lmm <- lmm(glucagonAUC ~ visit*group + (1|id/baseline), data = gastricbypassL)
eNRI.lmm
```

```
            Linear Mixed Model with nested random intercepts

outcome/cluster/time: glucagonAUC/id/XXtimeXX
data                 : 78 observations from 20 clusters
parameter            : 8 mean ((Intercept) visit2 visit3 visit4 group1 visit2:group1 visit3:group1 v
                       1 variance (sigma)
                       2 correlation (rho(id/baseline) rho(id))
log-restr.likelihood: -308.994835006264
convergence          : TRUE (6 iterations)
```

We obtain the same log-likelihood as, respectively, `eCS.lmm` and `eBCS.lmm`. Indeed, as previously mentioned, with positive residual correlation the random effect structure is equivalent to a compound symmetry structure.

⚠ random slopes are not currently supported in LMMstar.

⚠ the proposed implementation can be very inefficient compared to `lme4`.

The joint distribution between the outcome $\boldsymbol{Y}$ and the random effects $\boldsymbol{\eta}$ can be expressed as:

$$\begin{bmatrix} \boldsymbol{Y} \\ \boldsymbol{\eta} \end{bmatrix} \sim \mathcal{N}\left(\begin{bmatrix} \boldsymbol{\mu} \\ \boldsymbol{0} \end{bmatrix}, \begin{bmatrix} \Omega & Z\Psi \\ \Psi Z^{\mathsf{T}} & \Psi \end{bmatrix} \right)$$

Denote by $\varepsilon_i = \boldsymbol{Y}_i - \boldsymbol{\mu}_i$ the vector of marginal residuals relative to individual i, Ω_i its variance-covariance matrix, and $\psi_j = (\Psi)_{jj}$ the variance of the j-th random effect. We can re-express the expected value of the j-th random effect for individual i as:

$$\eta_{ij} = \psi_j Z_{ij} \Omega_i^{-1} \varepsilon_i$$

This is what the `ranef` method returns:

`head(ranef(eRI.lmm, format = "wide"))`	`head(ranef(eNRI.lmm, format = "wide"))`

```
  id estimate                    id   estimate estimate.FALSE estimate.TRUE
1  1 -2.51154                 1  1 -0.494271       -3.50959      -3.23209
2  2  1.01043                 2  2  0.186051      -10.39431      12.93198
3  3  6.08384                 3  3  1.088409        9.36327       5.48225
4  4 -6.62350                 4  4 -1.219596      -11.06703      -5.56784
5  5  0.39519                 5  5  0.081686       -0.71254       1.82672
6  6 -2.73384                 6  6 -0.503386       -7.81700       0.95098
```

It is also possible to extract the variance decomposition by setting the argument `effects` to `"variance"`:

`ranef(eRI.lmm, effects = "variance",` `format = "wide")`	`ranef(eNRI.lmm, effects = "variance",` `format = "wide")`

```
      type absolute relative              type absolute relative
1    total  380.580  1.00000      1    total  380.957 1.000000
2       id   82.741  0.21741      2       id   15.465 0.040595
3 residual  297.839  0.78259      3 baseline  210.938 0.553705
                                  4 residual  154.554 0.405700
```

Confidence intervals can also be obtained setting the argument `se` to `TRUE` and `format` equal to `"long"`:

`head(ranef(eRI.lmm, se = TRUE))`

```
  id estimate     se      df    lower   upper
1  1 -2.51154 2.3019 11.1302  -7.5708  2.5477
2  2  1.01043 2.1163 15.7355  -3.4821  5.5030
3  3  6.08384 2.9771  6.2085  -1.1421 13.3098
4  4 -6.62350 3.1114  5.8319 -14.2902  1.0432
5  5  0.39519 1.9661 23.8446  -3.6640  4.4543
6  6 -2.73384 2.2940 10.0189  -7.8438  2.3761
```

4.10 Sum of squares

⚠ The definition of the sum of squares is not straightforward with mixed models. Intuitively summing residuals across several outcomes will be hard to interpret unless all outcomes have the same variance. This is why LMMstar does not provide them. Nevertheless for specific covariance structure, namely independence and compound symmetry (with positive correlation) structure, sum of squares can be deduced from the lmm object - see appendix C for the theoretical derivations. Importantly, with these structures the residuals can be reparametrised as random effects plus independent residuals, i.e. $\Omega = Z\Psi Z^\intercal + \delta I$ where I is the identity matrix and δ the variance of these independent residuals.

Appendix C illustrate how to extract the sum of squares for univariate linear regression (i.e. independence structure) and here we illustrate the case of a compound symmetry structure. A key step is to extract from the lmm object the conditional residual variance δ:

```
sigma2 <- coef(eCS.lmm, effect = "variance")^2
tau <- coef(eCS.lmm, effect = "correlation")*sigma2
delta <- unname(sigma2 - tau)
```

This step will typically depend on the covariance structure. The residual sum of squares (SSE) equals the residual degrees of freedom times the conditional variance:

```
df.res <- df.residual(eCS.lmm)
SSE <- df.res * delta
c(df.res = df.res, SSE = SSE)
```

```
df.res    SSE
    70  20849
```

For the regression sum of squares (SSR), we first extract the mean parameters and their variance-covariance based on the expected information:

```
eBeta.lmm <- coef(eCS.lmm)
eVcov.lmm <- vcov(eCS.lmm, type.information = "expected")
```

Parameters are grouped with respect to the original variable:

```
attr(model.matrix(eCS.lmm),"assign")
```

```
[1] 0 1 1 1 2 3 3 3
```

So we respect this grouping when computing the normalized SSR:

```
SSRstar.time <- eBeta.lmm[2:4] %*% solve(eVcov.lmm[2:4,2:4]) %*% eBeta.lmm[2:4]
SSRstar.group <- eBeta.lmm[5] %*% solve(eVcov.lmm[5,5]) %*% eBeta.lmm[5]
```

The SSR is obtained by multiplying the normalized SSR by the conditional variance:

```
SSR.time <- as.double(SSRstar.time * delta)
SSR.group <- as.double(SSRstar.group * delta)
c(time = SSR.time, group = SSR.group)
```

```
   time    group
7872.19  643.57
```

4.11 Proportion of explained variance and partial correlation

For a univariate linear model with homoschedastic residual variance, the proportion of explained variance, also called partial R^2 or partial η^2, is defined as the ratio between sum of squares (e.g. Lakens (2013), equation 12):

$$R^2 = \frac{SSR}{SSR + SSE}$$

```
c(SSR.time/ (SSR.time + SSE),
  SSR.group/ (SSR.group + SSE))
```

```
[1] 0.274092 0.029944
```

Computing the SSR for each individual coefficients, taking its squared root, and multiplying by the sign of the corresponding coefficient leads to the partial correlation. This procedure extends to covariance structures that can be reparametrised as random effects plus independent residuals (see previous subsection) such as the compound symmetry with non-negative correlation.

⚠ for other covariance structures, especially when the variance may be repetition-dependent, the definition of explained variance/partial correlation is not straightforward.

```
eCS.R2 <- partialCor(eCS.lmm, R2 = TRUE)
summary(eCS.R2)
```

```
                 Partial correlation

             estimate    se   df  lower upper p.value
  visit2       -0.073 0.119 52.4 -0.311 0.165 0.54028
  visit3        0.438 0.089 51.4   0.26 0.616 < 1e-04
  visit4         0.07 0.119 52.4 -0.168 0.308 0.55876
  group1       -0.173 0.114 60.7 -0.402 0.056 0.13527
  visit2:group1 0.041 0.119 52.8 -0.198  0.28 0.73256
  visit3:group1 0.284 0.106   52  0.071 0.497 0.01007
  visit4:group1 0.314 0.103   52  0.107 0.521 0.00365
  ---------------------------------------------------
  Columns lower and upper contain 95% pointwise confidence intervals for each coefficient.
  Degrees of freedom were computed using a Satterthwaite approximation (column df).

                 Coefficient of determination (R2)

             estimate   se   df  lower upper p.value
  visit         0.274 0.08 50.5  0.114 0.434  0.0012
  group          0.03 0.04 60.7 -0.049 0.109  0.4520
  visit:group   0.147 0.073 51.7 <0.001 0.295  0.0500
  global        0.598 0.053 40.4  0.492 0.705  <1e-04
  ---------------------------------------------------
  Columns lower and upper contain 95% pointwise confidence intervals for each coefficient.
  Degrees of freedom were computed using a Satterthwaite approximation (column df).
```

Here the line "global" refer to the R2 for all covariates, computed based on the SSR relative to all mean parameters but the intercept.

⚠ `partialCor` will compute values for all types of mixed models. But their interpretation as partial correlation and proportion of explained variance outside the compound symmetry with non-negative correlation is questionnable.

<u>Note</u>: Other software packages like `effectsize::eta_squared` uses another formula to estimate the partial R2:

$$R^2 = \frac{Fdf_{num}}{Fdf_{num} + df_{denom}}$$

where F denote the F-statistic, df_{num} (resp. df_{denom}) the degrees of freedom of the numerator (resp. denominator) of this statistic. However since the calculation of degrees of freedom in LMM is approximate, I would expect this approach to be less reliable than the one of `partialCor` based on the SSR and SSE.

```
aCS.aov <- anova(eCS.lmm)$multivariate
setNames(with(aCS.aov, statistic*df.num/(statistic*df.num+df.denom)), aCS.aov$test)
```

```
     visit         group visit:group
  0.335374      0.033811    0.186290
```

4.12 Model diagnostic

The method `residuals` returns the residuals in the wide format:

```
eUN.diagW <- residuals(eUN.lmm, type = "normalized", format = "wide")
colnames(eUN.diagW) <- gsub("normalized.","",colnames(eUN.diagW))
head(eUN.diagW)
```

```
  id     r.-13      r.-1       r.1      r.13
1  1 -0.36029 -0.11344  0.377177 -1.45539
2  2  0.77339  2.12301 -0.232908 -0.10708
3  3  1.14219 -1.44778 -0.654876  2.01259
4  4 -0.77473  0.20612 -0.127117 -1.39519
5  5  0.22435        NA  0.011432 -0.15398
6  6  0.27439 -0.67308 -1.031131  0.42724
```

or in the long format:

```
eUN.diagL <- residuals(eUN.lmm, type = "normalized", format = "long", keep.data = TRUE)
head(eUN.diagL)
```

```
  id visit time weight glucagonAUC group baseline fitted r.normalized
1  1     1  -13  127.2      20.690     1     TRUE 25.904     -0.36029
2  2     1  -13  165.2      49.922     0     TRUE 38.729      0.77339
3  3     1  -13  109.7      42.434     1     TRUE 25.904      1.14219
4  4     1  -13  146.2      27.517     0     TRUE 38.729     -0.77473
5  5     1  -13  113.1      29.151     1     TRUE 25.904      0.22435
6  6     1  -13  158.8      42.700     0     TRUE 38.729      0.27439
```

Various type of residuals can be extract but the normalized one are recommanded when doing model checking. Diagnostic plots can then be generated by the user, or directly from the `lmm` object via the method `plot` (which internally calls the `residuals` method):

- misspecification of the mean structure

```
plot(eUN.lmm, type = "scatterplot")
```

- misspecification of the variance structure

```
plot(eUN.lmm, type = "scatterplot2")
```


- misspecification of the correlation structure

```
plot(eUN.lmm, type = "correlation", type.residual = "response")
plot(eUN.lmm, type = "correlation", type.residual = "normalized")
```


- residual distribution vs. normal distribution [3]:

```
plot(eUN.lmm, type = "qqplot", engine = "qqtest",
     facet = ~time, labeller = "label_both", facet_nrow=1)
## Note: the qqtest package to be installed to use the argument engine.plot = "qqtest"
```

---

[3]see Oldford (2016) for guidance about how to read quantile-quantile plots.

⚠ Deviation from the normal distribution does not necessarily question the validity of the statistical inference. Moreover, for variance and correlation parameters, normally distributed data is not enougth to ensure valid statistical inference. Instead one could assess whether the log-likelihood is locally quadratic as this ensures normally distributed estimates in finite samples (Geyer, 2013). Since the likelihood function is a multi-dimensional function this is not an easy task but one can look at specific 'slices' using the `profile` method:

```
eUN.lmm_profile <- profile(eUN.lmm, effects = c("sigma","rho(-13,-1)"))
plot(eUN.lmm_profile)
```


33

4.13 Visualize model fit

The fitted values can be displayed via the `plot` method:

```
## left panel
plot(eUN.lmm, type = "fit", color = "group", size.text = 20)
```

⚠ the shaded area represent 95% confidence intervals (CIs), i.e. is not adjusted for multiplicity over time. More explicit (but sometimes less readable) representation of the CIs can be obtained by setting the argument `ci.alpha` to `NA`:

```
## middle panel
plot(eUN.lmm, type = "fit", color = "group", ci.alpha = NA, size.text = 20)
```

It is also possible to display the observed values along with the fitted values by setting the argument `obs.alpha` to a strictly positive value below or equal to 1. This argument controls the transparency of the color used to display the observed values:

```
## right panel
plot(eUN.lmm, type = "fit", obs.alpha = 0.25, ci = FALSE, size.text = 20)
```


When considering continuous covariates, e.g.:

```
## add baseline weight
gastricbypassLB <- merge(gastricbypassL, gastricbypassW[c("id","weight1")], by = "id")

eUN.lmmB <- lmm(glucagonAUC ~ weight1 + visit*group, repetition = ~time|id,
                structure = "UN", data = gastricbypassLB)
```

The default graphical display can be confusing as it shows one curve per distinct set of covariate values, i.e. one line per subject:

```
## left panel
plot(eUN.lmmB, type = "fit", color = "group", ci = FALSE, size.text = 20)
```

It is possible to restrict the display specific to a covariate value via the argument `at`:

```
## middel panel
plot(eUN.lmmB, type = "fit", color = "group", ci = FALSE, size.text = 20,
     at = data.frame(weight1 = 150), obs.alpha = 0.2)
```

34

The `plot` method calls the `autoplot` methods which returns a list containing:

- a ggplot2 object (element `plot`)

- the dataset used to generate the ggplot2 object (element `data`)

This should ease further customization of the graphical display, e.g.:

```
## right panel
gg.traj <- autoplot(eUN.lmmB, type = "fit", color = "group", size.text = 20, facet =~id)
gg.traj$plot + theme(legend.position = "bottom")
```

4.14   Partial residuals

In a linear model where we split the covariates and mean parameters into two sets:

$$Y_i = X_{1,i}\beta_1 + X_{2,i}\beta_2 + \varepsilon_i$$

the partial residuals w.r.t. to the covariate(s) X_2 are defined by $\varepsilon_i^{X_2} = Y_i - X_{1,i}\beta_1$.
They can be computed via the `residuals` method:

```
df.pres <- residuals(eUN.lmmB, type = "partial", variable = "weight1", keep.data = TRUE)
head(df.pres)
```

```
  id visit time weight glucagonAUC group baseline weight1   fitted r.partial
1  1     1    1    -13   127.2       20.690     0    TRUE    127.2 -20.684  -25.3242
2  1     1    1    115.5              92.600     0   FALSE    127.2 -20.684  -12.2923
3  1     1    1     -1   120.7        20.535     0    TRUE    127.2 -20.684  -24.7703
4  1     1    1     13   108.1        43.434     0   FALSE    127.2 -20.684  -37.3259
5 10     1    1     13    90.9        57.942     0   FALSE    118.0 -19.188   -7.1423
6 10     1    1    1      99.3       103.728     0   FALSE    118.0 -19.188   11.7323
```

In the output, the X_1 covariates (`time` and `group`) have been set to the reference level (`-13` and `0`) for all observations. Confusion with the ordering variable from the `repetition` argument of `lmm` was avoided by using a different 'time' variable in the mean (`time`) and repetition argument (`visit`) when calling `lmm`. These residuals can be directly displayed via the `plot` method:

```
## left panel
plot(eUN.lmmB, type = "partial", variable = "weight1")
## right panel
plot(eUN.lmmB, type = "partial", variable = c("(Intercept)","weight1"))
```


The `plot` methods can handle one continuous and one categorical covariate (in addition to the intercept) to display interaction plots. In that case each observation/fitted line is colored according to the categorical covariate.

4.15 Statistical inference (single model, linear)

The `anova` method can be used to test one or several linear combinations of the model coefficients using Wald tests. By default, it will simultaneously test all parameters associated to a variable:

```
anova(eUN.lmm)
```

```
              Multivariate Wald test

                  F-statistic       df p.value
mean: visit            5.803 (3,16.9) 0.00647 **
    : group            3.926 (1,18.0) 0.06302  .
    : visit:group      2.762 (3,17.3) 0.07332  .
```

   Note that here the p-values are not adjust for multiple comparisons over variables. It is possible to specify a null hypothesis to be test: e.g. is there a change in average weight just after taking the treatment in the reference group:

```
anova(eUN.lmm, effects = c("visit3-visit2=0"))
```

```
              Multivariate Wald test


         F-statistic       df p.value
all: 1        14.318 (1,17.8) 0.00138 **
```

   One can also simulateneously tests several null hypotheses:

```
e.anova <- anova(eUN.lmm, effects = c("visit3-visit2=0","visit4-visit2=0"))
summary(e.anova)
```

```
               Multivariate Wald test


          F-statistic       df p.value
   all: 1       8.512 (2,17.2)  0.0027 **
   ----------------------------------
 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1.
 Degrees of freedom were computed using a Satterthwaite approximation (column df).

               Univariate Wald test

                   estimate    se   df  lower  upper p.value
   visit3 - visit2   36.167 9.558 17.8 13.381 58.953 0.00263 **
   visit4 - visit2    9.256 7.738   18 -9.192 27.704 0.38153
   ----------------------------------------------------------
 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1.
 Columns lower/upper/p.value adjusted for multiple comparisons -- max-test.
 (1e+05 samples have been used)
 Model-based standard errors are derived from the observed information (column se).
 Degrees of freedom were computed using a Satterthwaite approximation (column df).
```

or return all pairwise comparisons for a given factor using the `mcp` function of the multcomp package:

```
library(multcomp)
summary(anova(eUN.lmm, effects = mcp(visit = "Tukey")))
```

Singular contrast matrix: contrasts "3 - 2" "4 - 2" "4 - 3" have been removed.

                Multivariate Wald test

            F-statistic   df p.value
  all: visit   5.803 (3,16.9) 0.00647 **

  Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1.
  Degrees of freedom were computed using a Satterthwaite approximation (column df).

                Univariate Wald test

        estimate   se   df   lower upper p.value
  2 - 1   -4.734 2.776 17.5 -12.451 2.982 0.32482
  3 - 1   31.433 8.63 17.6   7.444 55.422 0.00860 **
  4 - 1   4.521 8.005   18 -17.731 26.774 0.93260
  3 - 2   36.167 9.558 17.8   9.597 62.737 0.00660 **
  4 - 2   9.256 7.738   18 -12.256 30.767 0.60663
  4 - 3 -26.912 7.448 16.4 -47.615 -6.209 0.00916 **

  Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1.
  Columns lower/upper/p.value adjusted for multiple comparisons -- max-test.
  (1e+05 samples have been used)
  Model-based standard errors are derived from the observed information (column se).
  Degrees of freedom were computed using a Satterthwaite approximation (column df).

Warning message:
In mcp2matrix(model, linfct = linfct) :
  covariate interactions found -- default contrast might be inappropriate

Here the `summary` method prints not only the global test but also the result associated to each hypothesis. The warning is triggered by the presence of an interaction between `visit` and `group`: the time effect is only tested here for the reference group. One should look also at the time effect in the other group before concluding about the possible absence of a time effect.

Special characters: special characters, such as parentheses or mathematical operators, can cause problems when using this formula-like interface to specify linear contrasts on parameters. This typically arises when testing (transformed) variance or correlation parameters, parentheses:

```
try(
  anova(eUN.lmm,
        effects = c("log(k).-1=0","log(k).1=0","log(k).13=0"))
)
```

```
Error in .anova_Wald(object, effects = effects, robust = robust, multivariate = multivariate,  :
  Possible mispecification of the argument 'effects' as running mulcomp::glht lead to the following
Error in parse(text = ex[i]) : <text>:1:7: unexpected symbol
1: log(k).
          ^
```

It is then advised to build a contrast matrix, e.g.:

```
name.coef <- rownames(confint(eUN.lmm, effects = "all"))
name.varcoef <- grep("^k",name.coef, value = TRUE)
C <- matrix(0, nrow = 3, ncol = length(name.coef), dimnames = list(name.varcoef, name.coef))
diag(C[name.varcoef,name.varcoef]) <- 1
C[,1:9]
```

	(Intercept)	visit2	visit3	visit4	group1	visit2:group1	visit3:group1	visit4:group1	sigma
k.-1	0	0	0	0	0	0	0	0	0
k.1	0	0	0	0	0	0	0	0	0
k.13	0	0	0	0	0	0	0	0	0

And then call the **anova** method specifying the null hypothesis via the contrast matrix:

```
anova(eUN.lmm, effects = C)
```

```
        Multivariate Wald test


     F-statistic        df p.value
all: 1      3.388 (3,25.7)  0.0332 *
```

4.16 Statistical inference (multiple models, linear)

It is possible to adjust for multiple testing across several linear contrasts that may originate from differente `lmm` using the approach of Pipper et al. (2012):

- fit the mixed models using `lmm`. The LMM must be fitted on the same dataset (or on subsets on a common larger dataset) with the same `repetition` argument.

- use the `anova` method to indicate which hypotheses are being tested

- combine the tests using `rbind`.

Here is an (artificial) example:

```
Manova <- rbind(anova(eInd.lmm, effects = "visit3:group1 = 0", robust = FALSE),
                anova(eCS.lmm, effects = "visit3:group1 = 0", robust = FALSE),
                anova(eUN.lmm, effects = "visit3:group1 = 0", robust = FALSE),
                name = c("Ind","CS","UN"))
summary(Manova)
```

```
              Multivariate Wald test

        Chi2-statistic       df p.value
  all: 1           116.9 (3,Inf)  <1e-04 ***
  ---------------------------------------
 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1.


              Univariate Wald test

                   estimate     se   df  lower  upper p.value
 Ind: visit3:group1   27.001 14.285 25.3 -1.482 55.485  0.0631 .
 CS: visit3:group1    27.481  11.09 52.8  5.369 49.593  0.0137 *
 UN: visit3:group1    27.571  12.42 17.8  2.808 52.335  0.0268 *
 -----------------------------------------------------------
 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1.
 Columns lower/upper/p.value adjusted for multiple comparisons -- max-test.
 (error when computing the adjusted columns lower/upper/p.value by numerical integration: 0.00073)
 Model-based standard errors are derived from the observed information (column se).
```

4.17 Statistical inference (single model, non-linear)

The `estimate` function can be used to test one or several non-linear combinations of model coefficients, using a first order delta method to quantify uncertainty. The combination has to be specified via a function (argument `f`). To illustrate its use consider an ANCOVA analysis:

$$Y_{i1} = \alpha + \beta Y_{i,0} + \gamma X_i + e_i$$

```
e.ANCOVA <- lm(weight4 ~ weight1 + group, data = gastricbypassW)
summary(e.ANCOVA)$coef
```

```
             Estimate Std. Error  t value    Pr(>|t|)
(Intercept) -5.92851   8.780064 -0.67522 5.0861e-01
weight1      0.82363   0.064116 12.84598 3.5247e-10
group        4.14046   2.533355  1.63438 1.2056e-01
```

We can replicate this analysis by first fitting a mixed model:

$$Y_{ij} = \alpha_j + \gamma_j X_i + \varepsilon_{i,j} \text{ where } \varepsilon_i \sim \mathcal{N}\left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} \sigma_1^2 & \rho\sigma_1\sigma_2 \\ \rho\sigma_1\sigma_2 & \sigma_2^2 \end{bmatrix}\right)$$

```
gastricbypassL14 <- gastricbypassL[gastricbypassL$visit %in% c(1,4),]
gastricbypassL14$visit <- droplevels(gastricbypassL14$visit)
e.lmmANCOVA <- lmm(weight ~ visit + visit:group, repetition = ~visit|id,
                   data = gastricbypassL14)
```

and then perform a first order delta-method:

```
lava::estimate(e.lmmANCOVA, f = function(p){
  c(Y1 = as.double(p["rho(1,4)"]*p["k.4"]),
    X1 = as.double(p["visit4:group1"]-p["rho(1,4)"]*p["k.4"]*p["visit1:group1"]))
})
```

```
    estimate       se      df    lower   upper    p.value
Y1   0.82363 0.062309  9.8746  0.68456 0.9627 1.3327e-07
X1   4.14046 2.461978 15.1613 -1.10227 9.3832 1.1309e-01
```

Indeed:

$$\mathbb{E}\left[Y_{i2}|Y_{i1}, X_i\right] = \alpha_2 + \gamma_2 X_i + \rho\frac{\sigma_2}{\sigma_1}\left(Y_{i1} - \alpha_1 - \gamma_1 X_i\right)$$

$$= \alpha_2 - \rho\frac{\sigma_2}{\sigma_1}\alpha_1 + \rho\frac{\sigma_2}{\sigma_1}Y_{i1} + \left(\gamma_2 - \rho\frac{\sigma_2}{\sigma_1}\gamma_1\right)X_i$$

We obtain identical estimate but different standard-errors/degrees of freedom compared to the univariate linear model approach. The later is to be prefer as it does not rely on approximation. The former is nevertheless useful as it can handle missing data in the outcome variable.

4.18 Baseline adjustment

In clinical trial the group and intervention variable often do not coincide, e.g., in presence of baseline measurement. In our running example, the first two measurement are pre-treatment (i.e. treatment should be "none") while the last two measurements are post-treatment (i.e. treatment should be 1 or 2). The baselineAdjustment function can be helpful to define a time varying treatment variable:

- where baseline takes a specific value

```
gastricbypassL$treat <- baselineAdjustment(gastricbypassL, variable = "group",
                           repetition = ~visit|id, constrain = c("1","2"),
                           new.level = "none")
table(treat = gastricbypassL$treat,
      visit = gastricbypassL$visit,
      group = gastricbypassL$group)
```

```
, , group = 0                              , , group = 1

        visit                                      visit
 treat    1  2  3  4                        treat    1  2  3  4
  none   10 10  0  0                         none   10 10  0  0
  0       0  0 10 10                         0       0  0  0  0
  1       0  0  0  0                         1       0  0 10 10
```

- where baseline corresponds to the reference group

```
gastricbypassL$treat2 <- baselineAdjustment(gastricbypassL, variable = "group",
                            repetition = ~visit|id, constrain = c("1","2"))
table(treat = gastricbypassL$treat2,
      visit = gastricbypassL$visit,
      group = gastricbypassL$group)
```

```
, , group = 0                              , , group = 1

        visit                                      visit
 treat  1  2  3  4                          treat  1  2  3  4
   0   10 10 10 10                            0   10 10  0  0
   1    0  0  0  0                            1    0  0 10 10
```

- including interactions with group

```
gastricbypassL$visitXtreat <- baselineAdjustment(gastricbypassL, variable = "group",
                                 repetition = ~visit|id, constrain = c("1","2"),
                                 collapse.time = ".")

table(treat = gastricbypassL$visitXtreat,
      visit = gastricbypassL$visit,
      group = gastricbypassL$group)
```

```
, , group = 0                                    , , group = 1

       visit                                           visit
   treat  1  2  3  4                             treat  1  2  3  4
     1   10  0  0  0                               1   10  0  0  0
     2    0 10  0  0                               2    0 10  0  0
    3.0   0  0 10  0                              3.0   0  0  0  0
    4.0   0  0  0 10                              4.0   0  0  0  0
    3.1   0  0  0  0                              3.1   0  0 10  0
    4.1   0  0  0  0                              4.1   0  0  0 10
```

We would then typically like to model group differences only after baseline (i.e. only at 1 week and 3 months after). This can be performed using the time varying treatment variable, e.g.:

```
eC.lmm <- lmm(glucagonAUC ~ visitXtreat, data = gastricbypassL,
              repetition = ~visit|id, structure = "UN")
coef(eC.lmm) ## change from baseline
```

```
(Intercept)    visitXtreat2 visitXtreat3.0 visitXtreat4.0 visitXtreat3.1 visitXtreat4.1
    32.3167         -2.7478        34.3703        11.6559        56.0581        27.6108
```

or

```
eC2.lmm <- lmm(glucagonAUC ~ 0 + visitXtreat, data = gastricbypassL,
               repetition = ~visit|id, structure = "UN")
coef(eC2.lmm) ## absolute value
```

```
visitXtreat1    visitXtreat2 visitXtreat3.0 visitXtreat4.0 visitXtreat3.1 visitXtreat4.1
    32.317          29.569         66.687         43.973         88.375         59.927
```

The parametrization however does not (directly) output treatment effects. Instead one may be tempted to use a formula like `treatment*time`. However this will lead to a non-indentifiable model. Indeed we are only able to estimate a total of 6 means when constraining the expected baseline value between the two groups to be the same. Therefore can at most identify 6 effects. However the design matrix for the interaction model:

```
colnames(model.matrix(glucagonAUC ~ treat*visit, data = gastricbypassL))
```

```
[1] "(Intercept)"   "treat0"        "treat1"        "visit2"        "visit3"        "visit4"
[7] "treat0:visit2" "treat1:visit2" "treat0:visit3" "treat1:visit3" "treat0:visit4" "treat1:visit4"
```

contains 12 parameters (i.e. 6 too many). Fortunately, the `lmm` will drop non-identifiable effects from the model and fit the resulting simplified model:

```
eC3.lmm <- lmm(glucagonAUC ~ treat2*visit, data = gastricbypassL,
               repetition = ~visit|id, structure = "UN")
```

```
Constant values in the design matrix for the mean structure.
Coefficients "treat21" "treat21:visit2" relative to interactions "treat2:visit" have been removed.
```

with the following coefficients:

```
model.tables(eC3.lmm)
```

```
              estimate       se      df    lower    upper    p.value
(Intercept)    32.3167   3.4764  19.003  25.0407  39.5927  1.6802e-08
visit2         -2.7478   1.9950  19.007  -6.9232   1.4276  1.8441e-01
visit3         34.3703   8.6111  15.161  16.0331  52.7076  1.1573e-03
visit4         11.6559   7.5601  17.328  -4.2715  27.5833  1.4119e-01
treat21:visit3 21.6878  12.2967  13.748  -4.7316  48.1072  9.9982e-02
treat21:visit4 15.9549   9.6174  12.374  -4.9296  36.8395  1.2224e-01
```

One can vizualize the baseline adjustment via the `plot` function:

```
plot(eC3.lmm, color = "group", ci = FALSE, size.text = 20, obs.alpha = 0.1)
```


and retrieve the treatment at each timepoint using the `effects` method:

```
effects(eC3.lmm, variable = "treat2", type = "difference")
```

```
            Difference in average counterfactual outcome
              w.r.t 'treat2' values

                 estimate       se    df   lower   upper  p.value
treat2=1-0(t=1)         0        0   Inf       0       0       NA
treat2=1-0(t=2)         0        0   Inf       0       0       NA
treat2=1-0(t=3)    21.688   12.297  13.7  -4.732  48.107    0.100 .
treat2=1-0(t=4)    15.955    9.617  12.4   -4.93   36.84    0.122
```

4.19 Predictions

Two types of predictions can be performed with the `predict` method:

- **static predictions** that are only conditional on the covariates:

```
news <- gastricbypassL[gastricbypassL$id==2,]
news$glucagon <- 0
predict(eUN.lmm, newdata = news, se = TRUE)
```

```
  estimate     se     df  lower  upper
1   38.729 4.5765 18.003 29.114 48.344
2   33.995 4.1002 17.897 25.377 42.612
3   70.162 8.6491 17.695 51.968 88.356
4   43.250 6.1883 18.005 30.249 56.251
```

which can be computing by creating a design matrix:

```
X.12 <- model.matrix(formula(eUN.lmm), news)
X.12
```

```
   (Intercept) visit2 visit3 visit4 group1 visit2:group1 visit3:group1 visit4:group1
2            1      0      0      0      0             0             0             0
22           1      1      0      0      0             0             0             0
42           1      0      1      0      0             0             0             0
62           1      0      0      1      0             0             0             0
attr(,"assign")
[1] 0 1 1 1 2 3 3 3
attr(,"contrasts")
attr(,"contrasts")$visit
[1] "contr.treatment"

attr(,"contrasts")$group
[1] "contr.treatment"
```

and then multiplying it with the regression coefficients:

```
X.12 %*% coef(eUN.lmm)
```

```
     [,1]
2  38.729
22 33.995
42 70.162
62 43.250
```

- **dynamic predictions** that are conditional on the covariates and the outcome measured at other timepoints. Consider two subjects for who we would like to predict the weight 1 week before the intervention based on the weight 3 months before the intervention:

```
newd <- rbind(
  data.frame(id = 1, time = -13, visit = "1", group = 0, glucagonAUC = coef(eUN.lmm)["(
    Intercept)"]),
  data.frame(id = 1, time = 1, visit = "3", group = 0, glucagonAUC = NA),
  data.frame(id = 2, time = -13, visit = "1", group = 0, glucagonAUC = 50),
  data.frame(id = 2, time = 1, visit = "3", group = 0, glucagonAUC = NA)
)
predict(eUN.lmm, newdata = newd, type = "dynamic", keep.data = TRUE)
```

```
  id time visit group glucagonAUC estimate     se     df  lower  upper
1  1  -13     1     0      38.729       NA     NA     NA     NA     NA
2  1    1     3     0          NA   70.162 8.3308 17.592 52.630 87.694
3  2  -13     1     0      50.000       NA     NA     NA     NA     NA
4  2    1     3     0          NA   75.888 9.6360 12.711 55.022 96.753
```

The first subjects starts with the average glucagon while the second starts with a much higher glucagon. The predicted glucagon after the operation for the first subject is then the average glucagon while it is predicted to be higher for the second subject due to the positive correlation over time. The predicted value is computed using the formula of the conditional mean for a Gaussian vector:

```
mu1 <- coef(eUN.lmm)["(Intercept)"]
mu3 <- mu1 + coef(eUN.lmm)["visit3"]
Omega_11 <- sigma(eUN.lmm)[1,1]
Omega_31 <- sigma(eUN.lmm)[3,1]
as.double(mu3 + Omega_31 * (50 - mu1) / Omega_11)
```

```
[1] 75.888
```

5 Equivalence with other statistical methods

5.1 Welch two sample t-test

A two sample t-test:

```
t.test(weight4 ~ group, data = gastricbypassW)
```

```
        Welch Two Sample t-test

data:  weight4 by group
t = 0.591, df = 17.7, p-value = 0.56
alternative hypothesis: true difference in means between group 0 and group 1 is not equal to 0
95 percent confidence interval:
 -11.736   20.916
sample estimates:
mean in group 0 mean in group 1
        104.66          100.07
```

is equivalent to an independent covariance pattern with a different variable for each group:

```
e.ttest4 <- lmm(weight4 ~ group, structure = IND(~group),
            data = gastricbypassW, trace = FALSE)
model.tables(e.ttest4)
```

```
            estimate    se      df    lower    upper    p.value
(Intercept)   104.66 5.1045  9.0018  93.113 116.207 7.2710e-09
group          -4.59 7.7607 17.6824 -20.916  11.736 5.6171e-01
```

5.2 Paired t-test

With complete data, a paired t-test:

```
t.test(gastricbypassW$weight4, gastricbypassW$weight1, paired = TRUE)
```

```
        Paired t-test

data:  gastricbypassW$weight4 and gastricbypassW$weight1
t = -17.2, df = 19, p-value = 5e-13
alternative hypothesis: true mean difference is not equal to 0
95 percent confidence interval:
 -29.848 -23.362
sample estimates:
mean difference
        -26.605
```

is equivalent to a LMM with an unstructured covariate pattern:

```
e.lmm2tt <- lmm(weight ~ visit, repetition = ~visit|id, structure = "UN",
                data = gastricbypassL)
model.tables(e.lmm2tt)["visit4",,drop=FALSE]
```

```
       estimate     se    df   lower   upper   p.value
visit4  -26.605 1.5494 18.964 -29.848 -23.362 5.1692e-13
```

5.3 Welch two sample t-test on the change

With complete data, a two sample t-test comparing the change from baseline:

```
gastricbypassW.0 <- gastricbypassW[gastricbypassW$group==0,]
gastricbypassW.1 <- gastricbypassW[gastricbypassW$group==1,]
t.test(gastricbypassW.0$weight4-gastricbypassW.0$weight1,
       gastricbypassW.1$weight4-gastricbypassW.1$weight1)
```

```
        Welch Two Sample t-test

data:  gastricbypassW.0$weight4 - gastricbypassW.0$weight1 and gastricbypassW.1$weight4 - gastricby
t = -2.11, df = 13, p-value = 0.055
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -12.16771    0.14771
sample estimates:
mean of x mean of y
   -29.61    -23.60
```

is equivalent to a LMM with a stratified unstructured covariate pattern:

```
e.lmm2tt2 <- lmm(weight ~ visit*group, repetition = ~visit|id, structure = UN(~group),
                 data = gastricbypassL)
model.tables(e.lmm2tt2)["visit4:group1",,drop=FALSE]
```

```
             estimate     se      df     lower   upper  p.value
visit4:group1     6.01 2.8511 13.009 -0.14908 12.169    0.055
```

5.4 Multiple Student's t-test

Multiple t-tests:

```
e.ttest1 <- lmm(weight1 ~ group, structure = IND(~group),
                data = gastricbypassW, trace = FALSE)
e.ttest2 <- lmm(weight2 ~ group, structure = IND(~group),
                data = gastricbypassW, trace = FALSE)
e.ttest3 <- lmm(weight3 ~ group, structure = IND(~group),
                data = gastricbypassW, trace = FALSE)
```

can be adjusted for multiple comparison by first using the `anova` function to specify the parameter of interest and combining the results using `rbind`:

```
e.mttest <- rbind(anova(e.ttest1, effects = "group=0"),
                  anova(e.ttest2, effects = "group=0"),
                  anova(e.ttest3, effects = "group=0"),
                  anova(e.ttest4, effects = "group=0"))
model.tables(e.mttest, method = "bonferroni")
```

```
               estimate     se     df   lower   upper p.value
weight1: group   -10.60 8.9717 17.965 -35.498 14.298       1
weight2: group    -9.50 8.3951 17.985 -32.795 13.795       1
weight3: group    -8.92 8.1295 17.959 -31.481 13.641       1
weight4: group    -4.59 7.7607 17.682 -26.165 16.985       1
```

⚠ efficient adjustment for multiple comparisons (like `"single-step"`) will not be valid as the correlation structure has not be specified. To do so it is more conveniently to work with a the long format:

```
e.mttest2 <- mlmm(weight ~ group, structure = IND(~group),
                  data = gastricbypassL, trace = FALSE,
                  effects = "group1=0", by = "time", repetition = ~time|id)
model.tables(e.mttest2, method = "single-step2")
```

```
   by parameter estimate     se     df   lower    upper p.value
1 -13    group1   -10.60 8.9717 17.965 -30.989   9.7893 0.31768
2  -1    group1    -9.50 8.3951 17.985 -28.579   9.5789 0.34123
3   1    group1    -8.92 8.1295 17.959 -27.395   9.5551 0.35785
4  13    group1    -4.59 7.7607 17.682 -22.227 13.0470 0.66473
```

or call the dedicated function `mt.test`:

```
mt.test(weight1+weight2+weight3+weight4~group, data = gastricbypassW)
```

```
        by parameter estimate     se     df   lower    upper p.value
1 weight1      group   -10.60 8.9717 17.965 -30.976   9.7758 0.31870
2 weight2      group    -9.50 8.3951 17.985 -28.566   9.5663 0.34115
3 weight3      group    -8.92 8.1295 17.959 -27.383   9.5429 0.35844
4 weight4      group    -4.59 7.7607 17.682 -22.215 13.0354 0.66272
```

5.5 Linear regression on the change

A widely spread approach to analyze longitudinal data is to reduce the number of repetitions to 1 by working on the change and then apply 'usual' statistical methods. For instance one could compare the pre- and post- operation values using:

```
gastricbypassW$changeG41 <- gastricbypassW$glucagonAUC4-gastricbypassW$glucagonAUC1
e.change41 <- lm(changeG41 ~ weight1, data = gastricbypassW)
summary(e.change41)$coef
```

```
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 88.41370   41.01024  2.1559 0.044871
weight1     -0.53331    0.31432 -1.6967 0.106975
```

This turns out to be equivalent to the following mixed model:

```
gastricbypassL41 <- gastricbypassL[gastricbypassL$visit %in% c(1,4),]
gastricbypassL41$visit <- droplevels(gastricbypassL41$visit)
gastricbypassL41$weight1 <- gastricbypassW$weight1[gastricbypassL41$id]

e.lmm41 <- lmm(glucagonAUC ~ visit + visit*weight1,
               repetition =~ visit|id, structure = "UN",
               data = gastricbypassL41)
model.tables(e.lmm41)
```

```
                estimate       se     df      lower      upper p.value
(Intercept)    31.7805917 23.58747 18.003 -17.77425   81.33543 0.19458
visit4         88.4137014 41.01024 18.001   2.25477  174.57264 0.04487
weight1         0.0041566  0.18078 18.003  -0.37565    0.38396 0.98191
visit4:weight1 -0.5333052  0.31432 18.001  -1.19366    0.12705 0.10697
```

This equivalence only holds as there is no missing data.

```
index.missing41 <- which(is.na(gastricbypassW$changeG41))
index.missing41
```

```
integer(0)
```

5.6 Correlation between changes

In some studies, one is interested in studying the relation between two evolutions. Say weight and glucagon before and after the operation:

```
gastricbypassW$changeG41 <- gastricbypassW$glucagonAUC4-gastricbypassW$glucagonAUC1
gastricbypassW$changeW41 <- gastricbypassW$weight4-gastricbypassW$weight1
```

One can evaluate their correlation:

```
cor.test(gastricbypassW$changeW41, gastricbypassW$changeG41)
```

```
        Pearson's product-moment correlation

data:  gastricbypassW$changeW41 and gastricbypassW$changeG41
t = 1.89, df = 18, p-value = 0.075
alternative hypothesis: true correlation is not equal to 0
95 percent confidence interval:
 -0.043829  0.719624
sample estimates:
    cor
0.40658
```

or regress one against the other:

```
e2.change41 <- lm(changeG41 ~ changeW41, data = gastricbypassW)
summary(e2.change41)$coef
```

```
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  65.0794   24.83368  2.6206 0.017331
changeW41     1.7082    0.90473  1.8881 0.075246
```

This problem can be recast using all measurement as outcomes:

```
keep.col <- c("id","weight1","weight4","glucagonAUC1","glucagonAUC4")
gastricbypassL4 <- reshape(gastricbypassW[,keep.col], direction = "long",
                        idvar = "id", varying = 2:5, timevar = "type", v.names = "value")
gastricbypassL4$type <- factor(gastricbypassL4$type, labels = keep.col[-1])
gastricbypassL4 <- gastricbypassL4[order(gastricbypassL4$id),]
head(gastricbypassL4)
```

```
    id        type   value
1.1  1      weight1 127.200
1.2  1      weight4 108.100
1.3  1 glucagonAUC1  20.690
1.4  1 glucagonAUC4  43.434
2.1  2      weight1 165.200
2.2  2      weight4 132.000
```

fitting an unstructured mixed model:

```
e.lmm4 <- lmm(value ~ type,
              repetition = ~type|id, structure = "UN",
              data = gastricbypassL4)
```

extract the residual covariance matrix:

```
sigma.lmm4 <- sigma(e.lmm4)
sigma.lmm4
```

	weight1	weight4	glucagonAUC1	glucagonAUC4
weight1	410.8475	326.84	1.7077	-217.399
weight4	326.8357	290.84	-24.6003	-161.696
glucagonAUC1	1.7077	-24.60	241.7007	-81.649
glucagonAUC4	-217.3994	-161.70	-81.6493	442.464

Deduce the residual covariance matrix for the change:

```
Mcon <- cbind(c(-1,1,0,0),c(0,0,-1,1))
sigmeChange.lmm4 <- t(Mcon) %*% sigma.lmm4 %*% Mcon
dimnames(sigmeChange.lmm4) <- list(c("d.weight","d.glucagonAUC"),
                                   c("d.weight","d.glucagonAUC"))
sigmeChange.lmm4
```

	d.weight	d.glucagonAUC
d.weight	48.011	82.011
d.glucagonAUC	82.011	847.464

and the corrrelation or covariance:

```
cov2cor(sigmeChange.lmm4)[1,2]
sigmeChange.lmm4[1,2]/sigmeChange.lmm4[1,1]
```

```
[1] 0.40658
[1] 1.7082
```

The uncertainty can be quantified using a delta method:

```
estimate(e.lmm4, function(p){
  Sigma.change <- t(Mcon) %*% sigma(e.lmm4, p = p) %*% Mcon
  c(cor = cov2cor(Sigma.change)[1,2],
    beta = Sigma.change[1,2]/Sigma.change[1,1])
})
```

	estimate	se	df	lower	upper	p.value
cor	0.40658	0.19150	2.5925	-0.26078	1.0739	0.13791
beta	1.70818	0.88073	2.6876	-1.28836	4.7047	0.15837

The standard errors and degrees of freedom do not match the univariate analysis, suggesting poor small sample properties of this technic.

6 Missing values and imputation

We reconsider the example of the previous section, but now in presence of missing values. The `summarize` function can be used to describe the amount of missing data at each repetition:

```
sss <- summarize(glucagonAUC ~ time, data = gastricbypassL, na.rm = TRUE)
cbind(sss[,1:4], pc = paste0(100 * sss$missing / (sss$missing + sss$observed), "%"))
```

```
       outcome time observed missing pc
1 glucagonAUC  -13       20       0 0%
2 glucagonAUC   -1       19       1 5%
3 glucagonAUC    1       19       1 5%
4 glucagonAUC   13       20       0 0%
```

For more detail about the missing data patters, see the `summarizeNA` function:

```
summarizeNA(data = gastricbypassL, repetition = ~ time|id)
```

variable	frequency	missing.pattern	n.missing	id	-13	-1	1	13
visit	20	00000	0	0	0	0	0	0
weight	20	00000	0	0	0	0	0	0
glucagonAUC	18	00000	0	0	0	0	0	0
	1	00010	1	0	0	0	1	0
	1	00100	1	0	0	1	0	0
group	20	00000	0	0	0	0	0	0
baseline	20	00000	0	0	0	0	0	0
treat	20	00000	0	0	0	0	0	0
treat2	20	00000	0	0	0	0	0	0
timeXtreat	20	00000	0	0	0	0	0	0
visitXtreat	20	00000	0	0	0	0	0	0
group2	20	00000	0	0	0	0	0	0

To begin with we will only consider 1 week before and 1 week after surgery:

```
## long format
gastricbypassL32 <- gastricbypassL[gastricbypassL$visit %in% c(3,2),]
gastricbypassL32$visit <- droplevels(gastricbypassL32$visit)
gastricbypassL32$weight1 <- gastricbypassW$weight1[gastricbypassL32$id]
## wide format
gastricbypassW$changeG32 <- gastricbypassW$glucagonAUC3-gastricbypassW$glucagonAUC2
```

6.1 Full information approach

LMM uses a full information approach:

```
e.lmm32 <- lmm(glucagonAUC ~ visit + visit*weight1,
               repetition =~ visit|id, structure = "UN",
               data = gastricbypassL32)
model.tables(e.lmm32)
```

```
                estimate       se      df      lower      upper    p.value
(Intercept)      9.24975 20.66248 17.011 -34.34202  52.84153 0.66004717
visit3         170.75489 39.66297 17.808  87.36171 254.14806 0.00043568
weight1          0.15750  0.15734 17.011  -0.17444   0.48945 0.33083625
visit3:weight1  -0.95238  0.30205 17.648  -1.58787  -0.31688 0.00560372
```

whereas a linear model would perform a complete case approach:

```
e.change32 <- lm(changeG32 ~ weight1, data = gastricbypassW)
summary(e.change32)$coef
```

```
              Estimate Std. Error t value    Pr(>|t|)
(Intercept) 173.46620    41.75201  4.1547 0.00074599
weight1      -0.96982     0.31589 -3.0701 0.00732363
```

In the former the likelihood is evaluated using all observations, even those from individuals with some (but not all) missing outcome values: baseline is used even if follow-up is missing. In the later the likelihood is only evaluated on individuals with no missing outcome values: if follow-up is missing then baseline is not used. Indeed:

```
coef(lm(changeG32 ~ weight1, data = gastricbypassW[-c(5,15),]))
```

```
(Intercept)       weight1
  173.46620      -0.96982
```

The estimates of the LMM can be retrived using a linear model where we have imputed the conditional expectation of the missing values given the observed value and the estimated model parameters: (see section 6.3 for a graphical representation)

```
gastricbypassWA <- fitted(e.lmm32, type = "outcome", format = "wide")
gastricbypassWA$change32 <- gastricbypassWA$glucagonAUC_3 - gastricbypassWA$glucagonAUC_2
gastricbypassWA$weight1 <- gastricbypassW$weight1[match(gastricbypassW$id,gastricbypassWA$id)]
coef(lm(change32 ~ weight1, data = gastricbypassWA))
```

```
(Intercept)       weight1
  170.75489      -0.95238
```

⚠ Standard errors, confidence intervals, and p-values from this linear model should not be trusted as they do not account for the uncertainty in the imputed values.

6.2 Complete case approach

The `lmmCC` can be used to obtain the LMM that is equivalent to a linear regression. In the case of the comparing the change between groups, the `repetition` argument should indicate how the change has been computed:

```
e.lmmCC <- lmmCC(e.change32, repetition = changeG32 ~ glucagonAUC3-glucagonAUC2|id)
model.tables(e.lmmCC)
```

```
Remove 2 clusters (4 observations)
 - 2 observations with missing data (2 clusters)
 - 0 missing repetitions (0 clusters)
              estimate       se df      lower      upper     p.value
(Intercept)  -165.90910 55.22956 16 -282.99061 -48.82760 0.00840925
time          173.46620 41.75201 16   84.95611 261.97630 0.00074594
weight1         1.13813  0.41786 16    0.25231   2.02395 0.01502328
time:weight1   -0.96982  0.31589 16   -1.63948  -0.30017 0.00732343
```

As output, the data from two clusters (i.e. 4 observations) has been excluded before fitting the LMM (instead of just the 2 observations with missing values for the full information approach). The interaction term of the LMM matches the regression coefficient of the linear model:

```
summary(e.change32)$coef
```

```
            Estimate Std. Error t value    Pr(>|t|)
(Intercept) 173.46620   41.75201  4.1547 0.00074599
weight1      -0.96982    0.31589 -3.0701 0.00732363
```

In the case of regressing two changes:

```
gastricbypassW$changeW32 <- gastricbypassW$weight3 - gastricbypassW$weight2

e2g.change32 <- lm(changeG32 ~ changeW32 + group, data = gastricbypassW)
summary(e2g.change32)$coef
```

```
            Estimate Std. Error  t value Pr(>|t|)
(Intercept)   16.998    30.9485  0.54924  0.59093
changeW32     -3.288     5.0127 -0.65594  0.52180
group         26.361    15.7415  1.67463  0.11472
```

the `repetition` argument should indicate how each change has been computed:

```
e2.lmmCC <-  lmmCC(e2g.change32, repetition = list(changeG32 ~ glucagonAUC3-glucagonAUC2|id,
                                                   changeW32 ~ weight3-weight2|id))
model.tables(e2.lmmCC)
```

```
Remove 2 clusters (8 observations)
 - 2 observations with missing data (2 clusters)
 - 0 missing repetitions (0 clusters)
      estimate      se     df    lower    upper p.value
cor   -0.16699 0.24255 1.7257  -1.3868   1.0529 0.57192
beta  -3.28804 4.83126 2.3131 -21.5987  15.0226 0.55791
```

We retrieve the same estimate for the effect of change in weights but the uncertainty (standard error, confidence intervals, p.value) do not match. They should be asymptotically correct but may not have very good small smaple properties.

6.3 Imputation

When fitting a linear mixed model on a dataset with missing values:

```
eUN.lmmNA <- lmm(glucagonAUC ~ time, repetition = ~time|id, data = gastricbypassL)
nobs(eUN.lmmNA)
```

```
obs          cluster      missing.obs missing.cluster
 78               20                2               0
```

It is possible to extract the most likely value for these missing observations using the `fitted` function with argument `impute=TRUE`:

```
eData <- fitted(eUN.lmmNA, type = "outcome", keep.data = TRUE)
eData <- eData[order(eData$id,eData$time),]
eData[eData$id %in% eData[eData$impute,"id"],c("id","visit","time","glucagonAUC","impute")]
```

```
   id visit time glucagonAUC impute
15 15     1  -13      22.244  FALSE
35 15     2   -1      32.544  FALSE
55 15     3    1      69.719   TRUE
75 15     4   13      85.222  FALSE
5   5     1  -13      29.151  FALSE
25  5     2   -1      26.270   TRUE
45  5     3    1      86.859  FALSE
65  5     4   13      57.970  FALSE
```

Missing outcome values in the dataset have been replaced by its most likely value (which is the same as the dynamic prediction, describedy previously). A column `impute` has also been added to differentiate between the the modeled and observed value. Visually:

```
ggplot(eData, aes(x=time,y=glucagonAUC, group=id)) + geom_line() + geom_point(aes(color=impute
    ))
```

It is possible to sample from the estimated distribution of the missing value instead of using the most likely value, e.g. accounting for residual variance and uncertainty related to parameter estimation:

```
index.na <- which(is.na(gastricbypassL$glucagonAUC))
set.seed(1)
fitted(eUN.lmmNA, type = "impute", se = c(TRUE,TRUE))[index.na]
set.seed(2)
fitted(eUN.lmmNA, type = "impute", se = c(TRUE,TRUE))[index.na]
set.seed(3)
fitted(eUN.lmmNA, type = "impute", se = c(TRUE,TRUE))[index.na]
```

```
[1] 21.932 75.390
[1] 20.060 75.428
[1] 19.610 60.684
```

6.4 Multiple imputation

The `mlmm` function can used to perform stratify analyses, typically useful when performing multiple imputations. Consider the wide format of the dataset where a few values are missing:

```
data(gastricbypassW, package = "LMMstar")
colSums(is.na(gastricbypassW))
```

id	weight1	weight2	weight3	weight4	glucagonAUC1	glucagonAUC2
0	0	0	0	0	0	1

glucagonAUC3	glucagonAUC4
1	0

We use `mice` to generate a number of imputed datasets (here 5):

```
library(mice)
set.seed(10)
gastricbypassW.mice <- mice(gastricbypassW, m = 5, printFlag = FALSE)
gastricbypassW.NNA <- complete(gastricbypassW.mice, action = "long")
table(gastricbypassW.NNA$.imp)
```

```
Warning message:
Number of logged events: 108
```

1	2	3	4	5
20	20	20	20	20

We can then use `mlmm` to perform a separate linear regression per dataset:

```
e.mlmm <- mlmm(glucagonAUC3~glucagonAUC2+weight2, data=gastricbypassW.NNA,
            by = ".imp", effects = "weight2=0", trace = FALSE)
model.tables(e.mlmm)
```

	by	parameter	estimate	se	df	lower	upper	p.value
1	1	weight2	-0.79725	0.31970	17.003	-1.4717	-0.12276	0.0232403
2	2	weight2	-0.83352	0.29798	17.003	-1.4622	-0.20486	0.0123742
3	3	weight2	-0.90658	0.28187	17.003	-1.5013	-0.31189	0.0050657
4	4	weight2	-0.90648	0.28183	17.003	-1.5011	-0.31189	0.0050638
5	5	weight2	-0.82477	0.30247	17.003	-1.4629	-0.18663	0.0143469

and pool the results using Rubin's rule:

```
model.tables(e.mlmm, method = "pool.rubin")
```

	estimate	se	df	lower	upper	p.value
<1, 5>	-0.85372	0.30212	14.741	-1.4987	-0.20878	0.012949

This matches (almost exactly, only the degrees of freedom are a little different) the results obtained with the mice package:

```
e.mice <- with(data=gastricbypassW.mice,exp=lm(glucagonAUC3~glucagonAUC2+weight2))
summary(pool(e.mice))
```

```
          term    estimate std.error  statistic     df    p.value
1  (Intercept) 178.988359  36.52589   4.900314 14.703 0.00020353
2 glucagonAUC2   0.027599   0.41848   0.065951 15.132 0.94828055
3      weight2  -0.853721   0.30212  -2.825775 14.737 0.01295073
```

One can use the `plot` function to obtain a forest plot of the individual estimates along with the pooled estimate:

```
plot(e.mlmm, method = c("pool.rubin","none"))
```

7 Data generation

Simulate some data in the wide format:

```
set.seed(10) ## ensure reproductibility
n.obs <- 100
n.times <- 4
mu <- rep(0,4)
gamma <- matrix(0, nrow = n.times, ncol = 10) ## add interaction
gamma[,6] <- c(0,1,1.5,1.5)
dW <- sampleRem(n.obs, n.times = n.times, mu = mu, gamma = gamma, format = "wide")
head(round(dW,3))
```

```
  id X1 X2 X3 X4 X5     X6     X7     X8    X9    X10     Y1     Y2     Y3     Y4
1  1  1  0  1  1  0 -0.367  1.534 -1.894 1.729  0.959  1.791  2.429  3.958  2.991
2  2  1  0  1  2  0 -0.410  2.065  1.766 0.761 -0.563  2.500  4.272  3.002  2.019
3  3  0  0  2  1  0 -1.720 -0.178  2.357 1.966  1.215 -3.208 -5.908 -4.277 -5.154
4  4  0  0  0  1  0  0.923 -2.089  0.233 1.307 -0.906 -2.062  0.397  1.757 -1.380
5  5  0  0  2  1  0  0.987  5.880  0.385 0.028  0.820  7.963  7.870  7.388  8.609
6  6  0  0  1  1  2 -1.075  0.479  2.202 0.900 -0.739  0.109 -1.602 -1.496 -1.841
```

Simulate some data in the long format:

```
set.seed(10) ## ensure reproductibility
dL <- sampleRem(n.obs, n.times = n.times, mu = mu, gamma = gamma, format = "long")
head(dL)
```

```
  id visit      Y X1 X2 X3 X4 X5       X6     X7      X8      X9     X10
1  1     1 1.7914  1  0  1  1  0 -0.36653 1.5338 -1.8944 1.72887 0.95925
2  1     2 2.4286  1  0  1  1  0 -0.36653 1.5338 -1.8944 1.72887 0.95925
3  1     3 3.9583  1  0  1  1  0 -0.36653 1.5338 -1.8944 1.72887 0.95925
4  1     4 2.9912  1  0  1  1  0 -0.36653 1.5338 -1.8944 1.72887 0.95925
5  2     1 2.5002  1  0  1  2  0 -0.40975 2.0654  1.7658 0.76133 -0.56302
6  2     2 4.2724  1  0  1  2  0 -0.40975 2.0654  1.7658 0.76133 -0.56302
```

8   Modifying default options

The `LMMstar.options` method enable to get and set the default options used by the package. For instance, the default option for the information matrix is:

```
LMMstar.options("type.information")
```

```
$type.information
[1] "observed"
```

To change the default option to "expected" (faster to compute but less accurate p-values and confidence intervals in small samples) use:

```
LMMstar.options(type.information = "expected")
```

To restore the original default options do:

```
LMMstar.options(reinitialise = TRUE)
```

9 R session

Details of the R session used to generate this document:

```
sessionInfo()
```

```
R version 4.3.3 (2024-02-29)
Platform: x86_64-pc-linux-gnu (64-bit)
Running under: Ubuntu 22.04.4 LTS

Matrix products: default
BLAS:   /usr/lib/x86_64-linux-gnu/blas/libblas.so.3.10.0
LAPACK: /usr/lib/x86_64-linux-gnu/lapack/liblapack.so.3.10.0

locale:
 [1] LC_CTYPE=en_US.UTF-8       LC_NUMERIC=C               LC_TIME=en_US.UTF-8
 [4] LC_COLLATE=en_US.UTF-8     LC_MONETARY=en_US.UTF-8    LC_MESSAGES=en_US.UTF-8
 [7] LC_PAPER=en_US.UTF-8       LC_NAME=C                  LC_ADDRESS=C
[10] LC_TELEPHONE=C             LC_MEASUREMENT=en_US.UTF-8 LC_IDENTIFICATION=C

time zone: Europe/Copenhagen
tzcode source: system (glibc)

attached base packages:
[1] grid      parallel  stats     graphics  grDevices utils     datasets  methods   base

other attached packages:
 [1] mice_3.16.0        emmeans_1.10.0    rlang_1.1.3       numDeriv_2016.8-1.1
 [5] doParallel_1.0.17  iterators_1.0.14  foreach_1.5.2     copula_1.1-3
 [9] multcomp_1.4-25    TH.data_1.1-2     MASS_7.3-60.0.1   survival_3.5-8
[13] mvtnorm_1.2-4      lme4_1.1-35.2     Matrix_1.6-5      lava_1.8.0
[17] nlme_3.1-163       LMMstar_1.1.0     ggpubr_0.6.0      ggplot2_3.5.1

loaded via a namespace (and not attached):
 [1] pbapply_1.7-2     gridExtra_2.3     pspline_1.0-19    remotes_2.5.0
 [5] sandwich_3.1-0    magrittr_2.0.3    butils.base_1.3   compiler_4.3.3
 [9] mgcv_1.9-1        systemfonts_1.0.6 vctrs_0.6.5       gsl_2.1-8
[13] stringr_1.5.1     profvis_0.3.8     shape_1.4.6.1     pkgconfig_2.0.3
[17] fastmap_1.1.1     backports_1.4.1   ellipsis_0.3.2    labeling_0.4.3
[21] utf8_1.2.4        promises_1.2.1    qqtest_1.2.0      sessioninfo_1.2.2
[25] nloptr_2.0.3      ragg_1.3.0        purrr_1.0.2       jomo_2.7-6
[29] glmnet_4.1-8      cachem_1.0.8      later_1.3.2       pan_1.9
[33] broom_1.0.5       R6_2.5.1          stringi_1.8.3     rpart_4.1.23
[37] parallelly_1.37.1 car_3.1-2         boot_1.3-30       pkgload_1.3.4
[41] estimability_1.5  Rcpp_1.0.12       future.apply_1.11.2 zoo_1.8-12
[45] usethis_2.2.3     nnet_7.3-19       httpuv_1.6.15     splines_4.3.3
[49] tidyselect_1.2.1  abind_1.4-5       codetools_0.2-19  miniUI_0.1.1.1
[53] listenv_0.9.1     pkgbuild_1.4.4    lattice_0.22-5    tibble_3.2.1
```

```
[57] shiny_1.8.1.1        withr_3.0.0         coda_0.19-4.1         future_1.33.2
[61] urlchecker_1.0.1     pillar_1.9.0        carData_3.0-5         stats4_4.3.3
[65] pcaPP_2.0-4          generics_0.1.3      munsell_0.5.1         scales_1.3.0
[69] minqa_1.2.6          globals_0.16.3      xtable_1.8-4          glue_1.7.0
[73] ADGofTest_0.3        tools_4.3.3         data.table_1.15.4     ggsignif_0.6.4
[77] fs_1.6.3             cowplot_1.1.3       tidyr_1.3.1           devtools_2.4.5
[81] colorspace_2.1-0     cli_3.6.2           textshaping_0.3.7     fansi_1.0.6
[85] dplyr_1.1.4          gtable_0.3.5        rstatix_0.7.2         stabledist_0.7-1
[89] digest_0.6.35        pbkrtest_0.5.2      htmlwidgets_1.6.4     farver_2.1.1
[93] memoise_2.0.1        htmltools_0.5.8.1   lifecycle_1.0.4       mitml_0.4-5
[97] mime_0.12
```

References

Christensen, R. (2011). *Plane answers to complex questions (4th edition)*, volume 35. Springer.

Geyer, C. J. (2013). Asymptotics of maximum likelihood without the lln or clt or sample size going to infinity. In *Advances in Modern Statistical Theory and Applications: A Festschrift in honor of Morris L. Eaton*, volume 10, pages 1–25. Institute of Mathematical Statistics.

Lakens, D. (2013). Calculating and reporting effect sizes to facilitate cumulative science: a practical primer for t-tests and anovas. *Frontiers in psychology*, 4:863.

Oldford, R. W. (2016). Self-calibrating quantile–quantile plots. *The American Statistician*, 70(1):74–90.

Pipper, C. B., Ritz, C., and Bisgaard, H. (2012). A versatile method for confirmatory evaluation of the effects of a covariate in multiple models. *Journal of the Royal Statistical Society: Series C (Applied Statistics)*, 61(2):315–326.

Appendix A   Likelihood in a linear mixed model

Denote by $\boldsymbol{Y}$ a vector of m outcomes, $\boldsymbol{X}$ a vector of p covariates, $\mu(\boldsymbol{\Theta}, \boldsymbol{X})$ the modeled mean, and $\Omega(\boldsymbol{\Theta}, \boldsymbol{X})$ the modeled residual variance-covariance. We consider n replicates (i.e. $\boldsymbol{Y}_1, \ldots, \boldsymbol{Y}_n$) and $V X_1, \ldots, \boldsymbol{X}_n$) along with a vector of weights $\omega = (w_1, \ldots, w_n)$, which are by default all equal to 1.

A.1   Log-likelihood

The restricted log-likelihood in a linear mixed model can then be written:

$$
\begin{aligned}
\mathcal{L}(\boldsymbol{\Theta}|\boldsymbol{Y}, \boldsymbol{X}) =& \frac{p}{2}\log(2\pi) - \frac{1}{2}\log\left(\left|\sum_{i=1}^{n} w_i \boldsymbol{X}_i \Omega_i^{-1}(\boldsymbol{\Theta})\boldsymbol{X}_i^{\intercal}\right|\right) \\
&+ \sum_{i=1}^{n} w_i\left(-\frac{m}{2}\log(2\pi) - \frac{1}{2}\log|\Omega_i(\boldsymbol{\Theta})| - \frac{1}{2}(\boldsymbol{Y}_i - \mu(\boldsymbol{\Theta}, \boldsymbol{X}_i))\Omega_i(\boldsymbol{\Theta})^{-1}(\boldsymbol{Y}_i - \mu(\boldsymbol{\Theta}, \boldsymbol{X}_i))^{\intercal}\right)
\end{aligned}
\tag{A}
$$

This is what the `logLik` method is computing for the REML criteria. The red term is specific to the REML criteria and prevents from computing individual contributions to the likelihood[4]. The blue term is what `logLik` outputs for the ML criteria when setting the argument `indiv` to `TRUE`.

A.2   Score

Using that $\partial\log(\det(X)) = tr(X^{-1}\partial(X))$, the score is obtained by derivating once the log-likelihood, i.e., for $\theta \in \boldsymbol{\Theta}$:

$$
\begin{aligned}
\mathcal{S}(\theta) =& \frac{\partial\mathcal{L}(\boldsymbol{\Theta}|\boldsymbol{Y}, \boldsymbol{X})}{\partial\theta} = \frac{1}{2}tr\left(\left(\sum_{i=1}^{n} w_i \boldsymbol{X}_i \Omega_i^{-1}(\boldsymbol{\Theta})\boldsymbol{X}_i^{\intercal}\right)^{-1}\left(\sum_{i=1}^{n} w_i \boldsymbol{X}_i \Omega_i^{-1}(\boldsymbol{\Theta})\frac{\partial\Omega_i(\boldsymbol{\Theta})}{\partial\theta}\Omega_i(\boldsymbol{\Theta})^{-1}\boldsymbol{X}_i^{\intercal}\right)\right) \\
&+ \sum_{i=1}^{n} w_i\left(-\frac{1}{2}tr\left(\Omega_i(\boldsymbol{\Theta})^{-1}\frac{\partial\Omega_i(\boldsymbol{\Theta})}{\partial\theta}\right) + \frac{\partial\mu(\boldsymbol{\Theta}, \boldsymbol{X}_i)}{\partial\theta}\Omega_i(\boldsymbol{\Theta})^{-1}(\boldsymbol{Y}_i - \mu(\boldsymbol{\Theta}, \boldsymbol{X}_i))^{\intercal}\right. \\
&\left.+ \frac{1}{2}(\boldsymbol{Y}_i - \mu(\boldsymbol{\Theta}, \boldsymbol{X}_i))\Omega_i(\boldsymbol{\Theta})^{-1}\frac{\partial\Omega_i(\boldsymbol{\Theta})}{\partial\theta}\Omega_i(\boldsymbol{\Theta})^{-1}(\boldsymbol{Y}_i - \mu(\boldsymbol{\Theta}, \boldsymbol{X}_i))^{\intercal}\right).
\end{aligned}
$$

This is what the `score` method is computing for the REML criteria. The red term is specific to the REML criteria and prevents from computing the score relative to each cluster. The blue term is what `score` outputs for the ML criteria when setting the argument `indiv` to `TRUE`.

---

[4]The REML is the likelihood of the observations divided by the prior on the estimated mean parameters $\widehat{\boldsymbol{\Theta}}_\mu \sim \mathcal{N}(\mu, (\boldsymbol{X}\Omega^{-1}(\boldsymbol{\Theta})\boldsymbol{X}^{\intercal})^{-1})$. This corresponds to $\frac{1}{\sqrt{2\pi}^p|(\sum_{i=1}^{n}\boldsymbol{X}_i\Omega_i^{-1}(\boldsymbol{\Theta})\boldsymbol{X}_i^{\intercal})^{-1}|}\exp\left(-(\widehat{\boldsymbol{\Theta}}_\mu - \mu)\left(2\sum_{i=1}^{n}\boldsymbol{X}_i\Omega_i^{-1}(\boldsymbol{\Theta})\boldsymbol{X}_i^{\intercal}\right)^{-1})(\widehat{\boldsymbol{\Theta}}_\mu - \mu)^{\intercal}\right)$ Since μ will be estimated to be $\boldsymbol{\Theta}_\mu$, the exponential term equals 1 and thus does not contribute to the log-likelihood. One divided by the other term gives $\sqrt{2\pi}^p\left(|\sum_{i=1}^{n}\boldsymbol{X}_i\Omega_i^{-1}(\boldsymbol{\Theta})\boldsymbol{X}_i^{\intercal}|\right)^{-1}$. The log of this term equals the red term

A.3 Hessian

Derivating a second time the log-likelihood gives the hessian, $\mathcal{H}(\mathbf{\Theta})$, with element[5]:

$$\mathcal{H}(\theta, \theta') = \frac{\partial^2 \mathcal{L}(\mathbf{\Theta}|\mathbf{Y}, \mathbf{X})}{\partial \theta \partial \theta'} = \frac{\partial \mathcal{S}(\theta)}{\partial \theta'}$$

$$= \frac{1}{2} tr \left(\left(\sum_{i=1}^{n} w_i \mathbf{X}_i \Omega_i^{-1}(\mathbf{\Theta}) \mathbf{X}_i^{\mathsf{T}} \right)^{-1} \left\{ \sum_{i=1}^{n} w_i \mathbf{X}_i \Omega_i^{-1}(\mathbf{\Theta}) \left(\frac{\partial^2 \Omega_i(\mathbf{\Theta})}{\partial \theta \partial \theta'} - 2 \frac{\partial \Omega_i(\mathbf{\Theta})}{\partial \theta} \Omega_i^{-1}(\mathbf{\Theta}) \frac{\partial \Omega_i(\mathbf{\Theta})}{\partial \theta'} \right) \Omega_i(\mathbf{\Theta})^{-1} \mathbf{X}_i^{\mathsf{T}} \right.$$

$$+ \left(\sum_{i=1}^{n} w_i \mathbf{X}_i \Omega_i^{-1}(\mathbf{\Theta}) \frac{\partial \Omega_i(\mathbf{\Theta})}{\partial \theta} \Omega_i(\mathbf{\Theta})^{-1} \mathbf{X}_i^{\mathsf{T}} \right) \left(\sum_{i=1}^{n} w_i \mathbf{X}_i \Omega_i^{-1}(\mathbf{\Theta}) \mathbf{X}_i^{\mathsf{T}} \right)^{-1} \left. \left(\sum_{i=1}^{n} w_i \mathbf{X}_i \Omega_i^{-1}(\mathbf{\Theta}) \frac{\partial \Omega_i(\mathbf{\Theta})}{\partial \theta'} \Omega_i(\mathbf{\Theta})^{-1} \mathbf{X}_i^{\mathsf{T}} \right) \right\} \right)$$

$$+ \sum_{i=1}^{n} w_i \left(\frac{1}{2} tr \left(\Omega_i(\mathbf{\Theta})^{-1} \frac{\partial \Omega_i(\mathbf{\Theta})}{\partial \theta'} \Omega_i(\mathbf{\Theta})^{-1} \frac{\partial \Omega_i(\mathbf{\Theta})}{\partial \theta} - \Omega_i(\mathbf{\Theta})^{-1} \frac{\partial^2 \Omega_i(\mathbf{\Theta})}{\partial \theta \partial \theta'} \right) \right.$$

$$- \frac{\partial \mu(\mathbf{\Theta}, \mathbf{X}_i)}{\partial \theta} \Omega_i(\mathbf{\Theta})^{-1} \frac{\partial \Omega_i(\mathbf{\Theta})}{\partial \theta'} \Omega_i(\mathbf{\Theta})^{-1} \varepsilon_i(\mathbf{\Theta})^{\mathsf{T}} - \frac{\partial \mu(\mathbf{\Theta}, \mathbf{X}_i)}{\partial \theta} \Omega_i(\mathbf{\Theta})^{-1} \frac{\partial \mu(\mathbf{\Theta}, \mathbf{X}_i)}{\partial \theta'}^{\mathsf{T}}$$

$$\left. + \frac{1}{2} \varepsilon_i(\mathbf{\Theta}) \Omega_i(\mathbf{\Theta})^{-1} \left(\frac{\partial^2 \Omega_i(\mathbf{\Theta})}{\partial \theta \partial \theta'} - \frac{\partial \Omega_i(\mathbf{\Theta})}{\partial \theta'} \Omega_i(\mathbf{\Theta})^{-1} \frac{\partial \Omega_i(\mathbf{\Theta})}{\partial \theta} - \frac{\partial \Omega_i(\mathbf{\Theta})}{\partial \theta} \Omega_i(\mathbf{\Theta})^{-1} \frac{\partial \Omega_i(\mathbf{\Theta})}{\partial \theta'} \right) \Omega_i(\mathbf{\Theta})^{-1} \varepsilon_i(\mathbf{\Theta})^{\mathsf{T}} \right).$$

where $\varepsilon_i(\mathbf{\Theta}) = \mathbf{Y}_i - \mu(\mathbf{\Theta}, \mathbf{X}_i)$.

The `information` method will (by default) return the (observed) information which is the opposite of the hessian. So multiplying the previous formula by -1 gives what `information` output for the REML criteria. The red term is specific to the REML criteria and prevents from computing the information relative to each cluster. The blue term is what `information` outputs for the ML criteria (up to a factor -1) when setting the argument `indiv` to `TRUE`.

A possible simplification is to use the expected hessian at the maximum likelihood. Indeed for any deterministic matrix A:

- $\mathbb{E}\left[A(\mathbf{Y}_i - \mu(\mathbf{\Theta}, \mathbf{X}_i))^{\mathsf{T}} | \mathbf{X}_i\right] = 0$

- $\mathbb{E}\left[(\mathbf{Y}_i - \mu(\mathbf{\Theta}, \mathbf{X}_i))A(\mathbf{Y}_i - \mu(\mathbf{\Theta}, \mathbf{X}_i))^{\mathsf{T}} || \mathbf{X}_i\right] = tr(A\mathbb{V}ar(\mathbf{Y}_i - \mu(\mathbf{\Theta}, \mathbf{X}_i)))$

when $\mathbb{E}\left[\mathbf{Y}_i - \mu(\mathbf{\Theta}, \mathbf{X}_i)\right] = 0$. This leads to:

$$\mathbb{E}\left[\mathcal{H}(\theta, \theta')|\mathbf{X}\right]$$

$$= \frac{1}{2} tr \left(\left(\sum_{i=1}^{n} w_i \mathbf{X}_i \Omega_i^{-1}(\mathbf{\Theta}) \mathbf{X}_i^{\mathsf{T}} \right)^{-1} \left\{ \sum_{i=1}^{n} w_i \mathbf{X}_i \Omega_i^{-1}(\mathbf{\Theta}) \left(\frac{\partial^2 \Omega_i(\mathbf{\Theta})}{\partial \theta \partial \theta'} - 2 \frac{\partial \Omega_i(\mathbf{\Theta})}{\partial \theta} \Omega_i^{-1}(\mathbf{\Theta}) \frac{\partial \Omega_i(\mathbf{\Theta})}{\partial \theta'} \right) \Omega_i(\mathbf{\Theta})^{-1} \mathbf{X}_i^{\mathsf{T}} \right.$$

$$+ \left(\sum_{i=1}^{n} w_i \mathbf{X}_i \Omega_i^{-1}(\mathbf{\Theta}) \frac{\partial \Omega_i(\mathbf{\Theta})}{\partial \theta} \Omega_i(\mathbf{\Theta})^{-1} \mathbf{X}_i^{\mathsf{T}} \right) \left(\sum_{i=1}^{n} w_i \mathbf{X}_i \Omega_i^{-1}(\mathbf{\Theta}) \mathbf{X}_i^{\mathsf{T}} \right)^{-1} \left. \left(\sum_{i=1}^{n} w_i \mathbf{X}_i \Omega_i^{-1}(\mathbf{\Theta}) \frac{\partial \Omega_i(\mathbf{\Theta})}{\partial \theta'} \Omega_i(\mathbf{\Theta})^{-1} \mathbf{X}_i^{\mathsf{T}} \right) \right\} \right)$$

$$+ \sum_{i=1}^{n} w_i \left(-\frac{1}{2} tr \left(\Omega_i(\mathbf{\Theta})^{-1} \frac{\partial \Omega_i(\mathbf{\Theta})}{\partial \theta'} \Omega_i(\mathbf{\Theta})^{-1} \frac{\partial \Omega_i(\mathbf{\Theta})}{\partial \theta} \right) - \frac{\partial \mu(\mathbf{\Theta}, \mathbf{X}_i)}{\partial \theta} \Omega_i(\mathbf{\Theta})^{-1} \frac{\partial \mu(\mathbf{\Theta}, \mathbf{X}_i)}{\partial \theta'}^{\mathsf{T}} \right) \quad \text{(B)}$$

This is what `information` output when the argument `type.information` is set to `"expected"` (up to a factor -1).

---

[5]if one is relative to the mean and the other to the variance then they are respectively θ and θ'

A.4 Degrees of freedom

Degrees of freedom are computed using a Satterthwaite approximation, i.e. for an estimate coefficient $\widehat{\beta} \in \widehat{\Theta}$ with standard error $\sigma_{\widehat{\beta}}$, the degree of freedom is:

$$df\left(\sigma_{\widehat{\beta}}\right) = \frac{2\sigma_{\widehat{\beta}}^4}{\mathbb{V}ar\left[\widehat{\sigma}_{\widehat{\beta}}\right]}$$

Using a first order Taylor expansion we can approximate the variance term as:

$$\mathbb{V}ar\left[\widehat{\sigma}_{\widehat{\beta}}\right] \approx \frac{\partial\widehat{\sigma}_{\widehat{\beta}}}{\partial\Theta}\Sigma_{\Theta}\frac{\partial\widehat{\sigma}_{\widehat{\beta}}}{\partial\Theta}^{\mathsf{T}}$$

$$\approx c_\beta\left(\widehat{\mathcal{I}}_{\widehat{\Theta}}\right)^{-1}\frac{\partial\widehat{\mathcal{I}}_{\widehat{\Theta}}}{\partial\Theta}\left(\widehat{\mathcal{I}}_{\widehat{\Theta}}\right)^{-1}c_\beta{}^{\mathsf{T}}\Sigma_{\Theta}c_\beta{}^{\mathsf{T}}\left(\widehat{\mathcal{I}}_{\widehat{\Theta}}\right)^{-1}\frac{\partial\widehat{\mathcal{I}}_{\widehat{\Theta}}}{\partial\Theta}^{\mathsf{T}}\left(\widehat{\mathcal{I}}_{\widehat{\Theta}}\right)^{-1}c_\beta$$

where Σ_{Θ} is the variance-covariance matrix of all model coefficients, $\mathcal{I}_{\Theta}$ the information matrix for all model coefficients, c_β a matrix used to select the element relative to β in the first derivative of the information matrix, and $\frac{\partial.}{\partial\Theta}$ denotes the vector of derivatives with respect to all model coefficients.

The derivative of the information matrix (i.e. negative hessian) can then be computed using numerical derivatives or using analytical formula. To obtain the later we first notice that:

$$\mathcal{H}(\theta, \theta') = \mathbb{E}\left[\mathcal{H}(\theta, \theta')|\boldsymbol{X}\right]$$
$$+ \sum_{i=1}^n w_i \left(tr\left(\Omega_i(\Theta)^{-1}\frac{\partial\Omega_i(\Theta)}{\partial\theta'}\Omega_i(\Theta)^{-1}\frac{\partial\Omega_i(\Theta)}{\partial\theta} - \Omega_i(\Theta)^{-1}\frac{\partial^2\Omega_i(\Theta)}{\partial\theta\partial\theta'}\right)\right.$$
$$- \frac{\partial\mu(\Theta, \boldsymbol{X}_i)}{\partial\theta}\Omega_i(\Theta)^{-1}\frac{\partial\Omega_i(\Theta)}{\partial\theta'}\Omega_i(\Theta)^{-1}\varepsilon_i(\Theta)^{\mathsf{T}}$$
$$\left.+ \frac{1}{2}\varepsilon_i(\Theta)\Omega_i(\Theta)^{-1}\left(\frac{\partial^2\Omega_i(\Theta)}{\partial\theta\partial\theta'} - \frac{\partial\Omega_i(\Theta)}{\partial\theta'}\Omega_i(\Theta)^{-1}\frac{\partial\Omega_i(\Theta)}{\partial\theta} - \frac{\partial\Omega_i(\Theta)}{\partial\theta}\Omega_i(\Theta)^{-1}\frac{\partial\Omega_i(\Theta)}{\partial\theta'}\right)\Omega_i(\Theta)^{-1}\varepsilon_i(\Theta)^{\mathsf{T}}\right)$$

$$\text{(C)}$$

where

$$\mathbb{E}\left[\mathcal{H}(\theta, \theta')|\boldsymbol{X}\right] = \frac{1}{2}tr\left(A(\Theta)^{-1}\left(\sum_{i=1}^n w_i b_i(\Theta)B_i(\Theta)b_i^{\mathsf{T}}(\Theta) + C(\Theta)A(\Theta)^{-1}C^{\mathsf{T}}(\Theta)\right)\right) + \sum_{i=1}^n w_i E_i(\Theta)$$

$$E_i(\Theta) = \frac{1}{2}tr\left(\Omega_i(\Theta)^{-1}\frac{\partial\Omega_i(\Theta)}{\partial\theta'}\Omega_i(\Theta)^{-1}\frac{\partial\Omega_i(\Theta)}{\partial\theta}\right) - \frac{\partial\mu(\Theta, \boldsymbol{X}_i)}{\partial\theta}\Omega_i(\Theta)^{-1}\frac{\partial\mu(\Theta, \boldsymbol{X}_i)}{\partial\theta'}^{\mathsf{T}}$$

$$A(\Theta) = \sum_{i=1}^n w_i\boldsymbol{X}_i\Omega_i^{-1}(\Theta)\boldsymbol{X}_i^{\mathsf{T}}$$

$$B(\Theta) = \frac{\partial^2\Omega_i(\Theta)}{\partial\theta\partial\theta'} - 2\frac{\partial\Omega_i(\Theta)}{\partial\theta}\Omega_i^{-1}(\Theta)\frac{\partial\Omega_i(\Theta)}{\partial\theta'}$$

$$b_i(\Theta) = \boldsymbol{X}_i\Omega_i^{-1}$$

$$C(\Theta) = \sum_{i=1}^n w_i\boldsymbol{X}_i\Omega_i^{-1}(\Theta)\frac{\partial\Omega_i(\Theta)}{\partial\theta}\Omega_i(\Theta)^{-1}\boldsymbol{X}_i^{\mathsf{T}}$$

So we will first derive the derivative of $\mathbb{E}\left[\mathcal{H}(\theta, \theta')|\boldsymbol{X}\right]$ and then the one of the blue term in Equation C. To simplify the derivation of the formula we will only derive them at the maximum likelihood, i.e. when

$\mathbb{E}\left[\frac{\partial \mathcal{H}(\theta,\theta'|\boldsymbol{X})}{\partial \theta''}\right] = \frac{\partial \mathbb{E}[\mathcal{H}(\theta,\theta'|\boldsymbol{X})]}{\partial \theta''}$ where the expectation is taken over $\boldsymbol{X}$. We first notice that the derivative with respect to the mean parameters is 0. So we just need to compute the derivative with respect to a variance parameter θ'':

$$\frac{\partial A(\boldsymbol{\Theta})^{-1}\left(\sum_{i=1}^{n} w_i b_i(\boldsymbol{\Theta}) B_i(\boldsymbol{\Theta}) b_i^{\mathsf{T}}(\boldsymbol{\Theta}) + C(\boldsymbol{\Theta}) A(\boldsymbol{\Theta})^{-1} C^{\mathsf{T}}(\boldsymbol{\Theta})\right)}{\partial \theta''}$$

$$=A(\boldsymbol{\Theta})^{-1}\frac{\partial A(\boldsymbol{\Theta})}{\partial \theta''}A(\boldsymbol{\Theta})^{-1}\left(\sum_{i=1}^{n} w_i b_i(\boldsymbol{\Theta}) B_i(\boldsymbol{\Theta}) b_i^{\mathsf{T}}(\boldsymbol{\Theta}) + C(\boldsymbol{\Theta}) A(\boldsymbol{\Theta})^{-1} C^{\mathsf{T}}(\boldsymbol{\Theta})\right)$$

$$+ A(\boldsymbol{\Theta})^{-1}\left(\sum_{i=1}^{n} w_i \left(\frac{\partial b_i(\boldsymbol{\Theta})}{\partial \theta''}B_i(\boldsymbol{\Theta})b_i^{\mathsf{T}}(\boldsymbol{\Theta}) + b_i(\boldsymbol{\Theta})\frac{\partial B_i(\boldsymbol{\Theta})}{\partial \theta''}b_i^{\mathsf{T}}(\boldsymbol{\Theta}) + b_i(\boldsymbol{\Theta})B_i(\boldsymbol{\Theta})\frac{\partial b_i^{\mathsf{T}}(\boldsymbol{\Theta})}{\partial \theta''}\right.\right.$$

$$\left.\left.+\ \frac{\partial C(\boldsymbol{\Theta})}{\partial \theta''}A^{-1}(\boldsymbol{\Theta})C^{\mathsf{T}}(\boldsymbol{\Theta}) + C(\boldsymbol{\Theta})A^{-1}\frac{\partial A(\boldsymbol{\Theta})}{\partial \theta''}A^{-1}C^{\mathsf{T}}(\boldsymbol{\Theta}) + C(\boldsymbol{\Theta})A^{-1}(\boldsymbol{\Theta})\frac{\partial C^{\mathsf{T}}(\boldsymbol{\Theta})}{\partial \theta''}\right)\right)$$

and

$$\frac{\partial E(\boldsymbol{\Theta})}{\partial \theta''} = \sum_{i=1}^{n} w_i \left(-\frac{1}{2}tr\left(-2\Omega_i(\boldsymbol{\Theta})^{-1}\frac{\partial \Omega_i(\boldsymbol{\Theta})}{\partial \theta''}\Omega_i(\boldsymbol{\Theta})^{-1}\frac{\partial \Omega_i(\boldsymbol{\Theta})}{\partial \theta'}\Omega_i(\boldsymbol{\Theta})^{-1}\frac{\partial \Omega_i(\boldsymbol{\Theta})}{\partial \theta}\right.\right.$$

$$\left.+\Omega_i(\boldsymbol{\Theta})^{-1}\frac{\partial^2 \Omega_i(\boldsymbol{\Theta})}{\partial \theta' \partial \theta''}\Omega_i(\boldsymbol{\Theta})^{-1}\frac{\partial \Omega_i(\boldsymbol{\Theta})}{\partial \theta} + \Omega_i(\boldsymbol{\Theta})^{-1}\frac{\partial \Omega_i(\boldsymbol{\Theta})}{\partial \theta'}\Omega_i(\boldsymbol{\Theta})^{-1}\frac{\partial^2 \Omega_i(\boldsymbol{\Theta})}{\partial \theta \partial \theta''}\right)$$

$$\left.+\frac{\partial \mu(\boldsymbol{\Theta}, \boldsymbol{X}_i)}{\partial \theta}\Omega_i(\boldsymbol{\Theta})^{-1}\frac{\partial \Omega_i(\boldsymbol{\Theta})}{\partial \theta''}\Omega_i(\boldsymbol{\Theta})^{-1}\frac{\partial \mu(\boldsymbol{\Theta}, \boldsymbol{X}_i)}{\partial \theta'}^{\mathsf{T}}\right)$$

where:

$$\frac{\partial A(\boldsymbol{\Theta})}{\partial \theta''} = \sum_{i=1}^{n} w_i \boldsymbol{X}_i \Omega_i^{-1}(\boldsymbol{\Theta})\frac{\partial \Omega_i(\boldsymbol{\Theta})}{\partial \theta''}\Omega_i^{-1}(\boldsymbol{\Theta})\boldsymbol{X}_i^{\mathsf{T}}$$

$$\frac{\partial b_i(\boldsymbol{\Theta})}{\partial \theta''} = \boldsymbol{X}_i \Omega_i^{-1}(\boldsymbol{\Theta})\frac{\partial \Omega_i(\boldsymbol{\Theta})}{\partial \theta''}\Omega_i^{-1}(\boldsymbol{\Theta})$$

$$\frac{\partial B_i(\boldsymbol{\Theta})}{\partial \theta''} = \frac{\partial^3 \Omega_i(\boldsymbol{\Theta})}{\theta \theta' \theta''}$$

$$-2\left(\frac{\partial^2 \Omega_i(\boldsymbol{\Theta})}{\partial \theta \partial \theta''}\Omega_i^{-1}(\boldsymbol{\Theta})\frac{\partial \Omega_i(\boldsymbol{\Theta})}{\partial \theta'} + \frac{\partial \Omega_i(\boldsymbol{\Theta})}{\partial \theta}\Omega_i^{-1}(\boldsymbol{\Theta})\frac{\partial \Omega_i(\boldsymbol{\Theta})}{\partial \theta''}\Omega_i^{-1}(\boldsymbol{\Theta})\frac{\partial \Omega_i(\boldsymbol{\Theta})}{\partial \theta'} + \frac{\partial \Omega_i(\boldsymbol{\Theta})}{\partial \theta}\Omega_i^{-1}(\boldsymbol{\Theta})\frac{\partial^2 \Omega_i(\boldsymbol{\Theta})}{\partial \theta' \partial \theta''}\right)$$

$$\frac{\partial C(\boldsymbol{\Theta})}{\partial \theta''} = \sum_{i=1}^{n} w_i \boldsymbol{X}_i \Omega_i^{-1}(\boldsymbol{\Theta})\left(\frac{\partial \Omega_i(\boldsymbol{\Theta})}{\partial \theta''}\Omega_i^{-1}(\boldsymbol{\Theta})\frac{\partial \Omega_i(\boldsymbol{\Theta})}{\partial \theta} + \frac{\partial^2 \Omega_i(\boldsymbol{\Theta})}{\partial \theta \partial \theta''} + \frac{\partial \Omega_i(\boldsymbol{\Theta})}{\partial \theta}\Omega_i^{-1}(\boldsymbol{\Theta})\frac{\partial \Omega_i(\boldsymbol{\Theta})}{\partial \theta''}\right)\Omega_i^{-1}(\boldsymbol{\Theta})\boldsymbol{X}_i^{\mathsf{T}}$$

Appendix B   Likelihood ratio test with the REML criterion

The blue term of Equation A in the log-likelihood is invariant to re-parameterisation while the red term is not. This means that a re-parametrisation of X into $\tilde{X} = BX$ with B invertible would not change the likelihood when using ML but would decrease the log-likelihood by $\log(|B|)$ when using REML.

Let's take an example:

```
## data(dfL, package = "LMMstar")
dfTest <- gastricbypassL[!is.na(gastricbypassL$glucagonAUC),]
dfTest$gluc <- dfTest$glucagonAUC
dfTest$gluc2 <- dfTest$glucagonAUC*2
```

where we multiply one column of the design matrix by 2. As mentionned previously this does not affect the log-likelihood when using ML:

```
eML.UN <- lmm(weight ~ time+gluc, data = dfTest, repetition = ~time|id, method = "ML")
eML.UN2 <- lmm(weight ~ time+gluc, data = dfTest, repetition = ~time|id, method = "ML")
c(logLik(eML.UN), logLik(eML.UN2), logLik(eML.UN) - logLik(eML.UN2))
```

```
[1] -230.62 -230.62     0.00
```

but it does when using REML:

```
eREML.UN <- lmm(weight ~ time + gluc, data = dfTest, repetition = ~time|id, method = "REML")
eREML.UN2 <- lmm(weight ~ time + gluc2, data = dfTest, repetition = ~time|id, method = "REML"
    )
c(logLik(eREML.UN), logLik(eREML.UN2), logLik(eREML.UN) - logLik(eREML.UN2), log(2))
```

```
[1] -235.23462 -235.92777     0.69315     0.69315
```

Therefore, when comparing models with different mean effects there is a risk that the difference (or part of it) in log-likelihood is due to a new parametrisation and no only to a difference in model fit. This would typically be the case when adding an interaction where we can have a smaller restricted log-likehood when considering a more complex model:

```
set.seed(5)
dfTest$ff <- rbinom(NROW(dfTest), size = 1, prob = 0.5)
logLik(lmm(weight ~ time+gluc, data = dfTest, repetition = ~time|id, method = "REML"))
logLik(lmm(weight ~ time+gluc*ff, data = dfTest, repetition = ~time|id, method = "REML"))
```

```
[1] -235.23
[1] -238.93
```

This is quite counter-intuitive as more complex model should lead to better fit and would never happen when using ML:

```
logLik(lmm(weight ~ time + gluc, data = dfTest, repetition = ~time|id, method = "ML"))
logLik(lmm(weight ~ time + gluc*ff, data = dfTest, repetition = ~time|id, method = "ML"))
```

```
[1] -230.62
[1] -230.44
```

This is why, unless one knows what he/she is doing, it is not recommended to use likelihood ratio test to assess relevance of mean parameters in mixed models estimated with REML.

Appendix C Sum of squares in a linear mixed model

All mixed models implemented in LMMstar can be written as:

$$Y_{it} = X_{it}\beta + \varepsilon_{it} \text{ where } \varepsilon_i \sim \mathcal{N}(0, \Omega)$$

where Y denote the outcome repeatedly measured within each cluster i where t indexes the repetitions. X denotes the covariates, β the mean parameters, ε the residuals, and Ω the residual variance-covariance matrix. Ω must be positive definite so there must exist a square positive definite matrix $\Omega^{1/2}$ such that $\Omega^{1/2}\Omega^{1/2} = \Omega$. Therefore the previous model is equivalent to:

$$Y_{it}^* = X_{it}^*\beta + \varepsilon_{it}^* \text{ where } \varepsilon_i \sim \mathcal{N}(0, I_T)$$

where $Y_i^* = \Omega^{-1/2}Y_i$, $X_i^* = \Omega^{-1/2}X_i$, $\varepsilon_i^* = \Omega^{-1/2}\varepsilon_i$, and I_x is the identity matrix with x rows and columns. One can then introduce the projectors $H = X(X^{\intercal}\Omega^{-1}X)^{-1}X^{\intercal}\Omega^{-1}$ and $H^* = X^*(X^{*\intercal}X^*)^{-1}X^{*\intercal}$ onto the space spanned by X and X^* respectively. We can now define the "normalized" residual sum of squares as the squared sum of the normalized residuals:

$$\begin{aligned}
SSE^* = \varepsilon^{*\intercal}\varepsilon^* &= Y^{*\intercal}(I_{nT} - H^*)Y^* \\
&= Y^{\intercal}\Omega^{-1}Y - Y^{\intercal}\Omega^{-1}X\left(X^{\intercal}\Omega^{-1}X\right)^{-1}X^{\intercal}\Omega^{-1}Y \\
&= Y^{\intercal}(I_{nT} - H^{\intercal})\Omega^{-1}(I_{nT} - H)Y
\end{aligned}$$

The previous to last line uses that: $(I_{nT} - H^{\intercal})\Omega^{-1}(I_{nT} - H) = \Omega^{-1} - H^{\intercal}\Omega^{-1} - \Omega^{-1}H + H^{\intercal}\Omega^{-1}H = \Omega^{-1} - H^{\intercal}\Omega^{-1}$ as $H^{\intercal}\Omega^{-1}H = \Omega^{-1}HH = \Omega^{-1}H$ since H is a projector. Note that compared to the "traditional" SSE defined for linear regression and random effect models (e.g. see Christensen (2011) section 2.7), $SSE = \delta SSE^*$ where δ is the residual variance conditional on any random effects, i.e. SSE^* are the residual degrees of freedom. This is because the same definition for the sum of squares is used except that $\varepsilon_i \sim \mathcal{N}(0, \delta\Omega)$.

We can also define the "normalized" regression sum of squares:

$$\begin{aligned}
SSR^* = (X^*\beta)^{\intercal}X^*\beta &= (H^*Y^*)^{\intercal}H^*Y^* = Y^{*\intercal}H^*Y^* \\
&= Y^{\intercal}H^{\intercal}\Omega^{-1}Y^* = Y^{\intercal}H^{\intercal}H^{\intercal}\Omega^{-1}Y^* = Y^{\intercal}H^{\intercal}\Omega^{-1}HY^* \\
&= \widehat{\beta}X^{\intercal}\Omega^{-1}X\widehat{\beta}
\end{aligned}$$

where $\widehat{\beta} = (X^{\intercal}\Omega^{-1}X)^{-1}X^{\intercal}\Omega^{-1}Y$. Note that when using the expected information $SSR^* = \widehat{\beta}\Sigma_{\widehat{\beta}}^{-1}\widehat{\beta}$, i.e. it is the F-statistics times the number of parameters. Again the "traditional" SSR defined for linear regression and random effect models is proportional to this normalized SSR: $SSR = \delta SSR^*$.

The proportion of explained variance of p parameters can thus be re-expressed as:

$$R^2 = \frac{SSR}{SSR + SSE} = \frac{SSR^*}{SSR^* + SSE^*} = \frac{Fp}{Fp + df}$$

where df denotes the residual degrees of freedom, typically $n - p$ in a univariate linear model fitted with n observations.

⚠ In practice *df* is estimated using the Satterthwaite approximation of the degrees of freedom of the regression coefficient. This is only equivalent to the "SSR/SSE" formula in univariate linear regression.

Illustration for a univariate linear model:

Data without missing values:

```
df.aov <- gastricbypassL[!is.na(gastricbypassL$glucagon),]
```

Traditional anova decomposition:

```
e.lm <- lm(weight ~ visit + glucagonAUC, data = df.aov)
car::Anova(e.lm, type = "II")
```

```
Anova Table (Type II tests)

Response: weight
            Sum Sq Df F value Pr(>F)
visit         5837  3    5.94 0.0011 **
glucagonAUC   2133  1    6.51 0.0128 *
Residuals    23925 73
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Fit `lmm`:

```
e.lmm <- lmm(weight ~ visit + glucagonAUC, data = df.aov)
```

Residual sum of squares (SSE):

```
SSEstar <- crossprod(residuals(e.lmm, type = "normalized"))
c(SSEstar = SSEstar, SSE = SSEstar * sigma(e.lmm))
```

```
SSEstar     SSE
     73   23925
```

The normalized SSE can also be obtained using the `df.residual` method:

```
df.residual(e.lmm)
```

```
[1] 73
```

Regression sum of squares (SSR):

```
eBeta.lmm <- coef(e.lmm)
eVcov.lmm <- vcov(e.lmm, type.information = "expected")

SSRstar.glucagon <- eBeta.lmm[5] %*% solve(eVcov.lmm[5,5]) %*% eBeta.lmm[5]
SSRstar.time <- eBeta.lmm[2:4] %*% solve(eVcov.lmm[2:4,2:4]) %*% eBeta.lmm[2:4]
c(SSR.glucagon = SSRstar.glucagon * sigma(e.lmm),
  SSR.time = SSRstar.time * sigma(e.lmm),
  F.glucagon = SSRstar.glucagon,
  F.time = SSRstar.time/3)
```

```
SSR.glucagon     SSR.time   F.glucagon       F.time
   2132.629      5837.410        6.507        5.937
```

So the proportion of explained variance is:

```
R2.glucagon <- SSRstar.glucagon/(SSRstar.glucagon+SSEstar)
R2.glucagon
```

```
          [,1]
[1,] 0.081842
```

and the corresponding partial correlation is:

```
sign(coef(e.lmm)["glucagonAUC"])*sqrt(R2.glucagon)
```

```
          [,1]
[1,] -0.28608
```

which matches the output of `partialCor`:

```
summary(partialCor(e.lmm, R2 = TRUE))
```

```
                 Partial correlation

             estimate    se df  lower   upper p.value
   visit2       -0.151 0.113 73 -0.377   0.074 0.18450
   visit3       -0.013 0.117 73 -0.246    0.22 0.91230
   visit4       -0.381 0.092 73 -0.565  -0.197 < 1e-04
   glucagonAUC  -0.286 0.103 73 -0.491  -0.081 0.00695
   -------------------------------------------------
 Columns lower and upper contain 95% pointwise confidence intervals for each coefficient.
 Degrees of freedom were computed using a Satterthwaite approximation (column df).

             Coefficient of determination (R2)

             estimate    se df  lower upper  p.value
   visit        0.196 0.075 73  0.047 0.345 0.010548
   glucagonAUC  0.082 0.059 73 -0.036 0.199 0.169016
   global        0.29 0.075 73   0.14  0.44 0.000257
   -------------------------------------------------
 Columns lower and upper contain 95% pointwise confidence intervals for each coefficient.
 Degrees of freedom were computed using a Satterthwaite approximation (column df).
```

Appendix D Equivalence with other R packages

D.1 nlme package

The model class obtained with the `lmm` function overlaps the model class of the `lme` and `gls` functions from the nlme package.

```
library(nlme)
```

For instance, the compound symmetry is equivalent to `corCompSymm` correlation structure, or to a random intercept model (when the within subject correlation is positive):

```
eRI.lmm <- lmm(weight ~ visit*group, structure = "RE",
               data = gastricbypassL, repetition = ~visit|id)
eCS.gls <- gls(weight ~ visit*group, correlation = corCompSymm(form=~visit|id),
               data = gastricbypassL, na.action = na.omit)
eCS.lme <- lme(weight ~ visit*group, random = ~1|id,
               data = gastricbypassL, na.action = na.omit)
logLik(eRI.lmm)
logLik(eCS.lme)
logLik(eCS.gls)
```

```
[1] -236.21
'log Lik.' -236.21 (df=10)
'log Lik.' -236.21 (df=10)
```

The estimated random effect also match:

```
range(ranef(eRI.lmm)-ranef(eCS.lme))
```

```
[1] -1.7303e-08  2.6979e-08
```

Unstructured residual covariance matrix can also be obtained with `gls`:

```
eUN.gls <- gls(glucagonAUC ~ visit*group,
               correlation = corSymm(form=~as.numeric(visit)|id),
               weights = varIdent(form=~1|visit),
               data = gastricbypassL, na.action = na.omit)
logLik(eUN.gls)
logLik(eUN.lmm)
```

```
'log Lik.' -295.31 (df=18)
[1] -295.31
```

D.2 lme4 package

The model class obtained with the `lmm` function overlaps the model class of the `lmer` function from the lme4 package.

```
library(lme4)
library(lmerTest)
```

For instance, the compound symmetry is equivalent to a random intercept model (when the within subject correlation is positive):

```
eRI.lmer <- lmer(weight ~ visit*group + (1|id),
                 data = gastricbypassL)
logLik(eRI.lmer)
logLik(eRI.lmm)
```

```
'log Lik.' -236.21 (df=10)
[1] -236.21
```

The estimated random effects match:

```
range(ranef(eRI.lmm)-ranef(eRI.lmer)$id)
```

```
[1] -1.5513e-08  2.4171e-08
```

Nested random effects correspond to block unstructured:

```
eNRI.lmm <- lmm(weight ~ visit*group, structure = RE(~(1|id/baseline)),
               data = gastricbypassL, repetition = ~visit|id)
eNRI.lmer <- lmer(weight ~ visit*group + (1|id/baseline),
                  data = gastricbypassL)
logLik(eNRI.lmer)
logLik(eNRI.lmm)
```

```
'log Lik.' -234.97 (df=11)
[1] -234.97
```

And the estimated random effects still match:

```
eRanefNRI.lmm <- ranef(eNRI.lmm, format = "wide")
eRanefNRI.lmer <- ranef(eNRI.lmer)
## id
range(eRanefNRI.lmm$estimate-eRanefNRI.lmer$id)
## baseline
range(c(eRanefNRI.lmm$estimate.FALSE,eRanefNRI.lmm$estimate.TRUE)-ranef(eNRI.lmer)$`baseline:
    id`)
```

```
[1] -5.8317e-06  9.0913e-06
[1] -8.5850e-05  7.8971e-05
```

An unstructure residual covariance matrix can also be obtained using random slopes:

```
eUN.lmer <- lmer(glucagonAUC ~ visit*group + (0 + visit|id),
                 data = gastricbypassL,
                 control = lmerControl(check.nobs.vs.nRE = "ignore"))
logLik(eUN.lmer)
logLik(eUN.lmm)
```

```
Warning message:
In checkConv(attr(opt, "derivs"), opt$par, ctrl = control$checkConv,  :
  Model failed to converge with max|grad| = 0.00203036 (tol = 0.002, component 1)
'log Lik.' -295.31 (df=19)
[1] -295.31
```

The uncertainty is quantified in a slightly different way, e.g.:

```
anova(eUN.lmm)
```

```
        Multivariate Wald test

                F-statistic      df p.value
mean: visit           5.803 (3,16.9) 0.00647 **
    : group           3.926 (1,18.0) 0.06302  .
    : visit:group     2.762 (3,17.3) 0.07332  .
```

is very similar but not identical to:

```
## only the last line is comparable
anova(eUN.lmer)
```

```
Type III Analysis of Variance Table with Satterthwaite's method
            Sum Sq Mean Sq NumDF DenDF F value  Pr(>F)
visit         1339     446     3  17.4   18.29 1.3e-05 ***
group            5       5     1  18.1    0.22   0.647
visit:group    203      68     3  17.4    2.77   0.073 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

It is also possible to fit cross-random effects such as:

```
data("Penicillin")
eCRI.lmer <- lmer(diameter ~ 1 + (1|plate) + (1|sample), Penicillin)
logLik(eCRI.lmer)
```

```
'log Lik.' -165.43 (df=4)
```

using lmm:

```
Penicillin$index <- paste(Penicillin$sample,Penicillin$plate,sep=".")
Penicillin$id <- 1


eCRI.lmm <- lmm(diameter ~ 1 + (1|plate) + (1|sample), data = Penicillin)
logLik(eCRI.lmm)
```

[1] -165.43

Despite being significantly slower, the loglikelihood and random effect still match:

```
range(ranef(eCRI.lmm)$estimate-rbind(ranef(eCRI.lmer)$plate,ranef(eCRI.lmer)$sample))
```

[1] -4.3812e-07 6.0172e-07

D.3 mmrm package

The package `mmrm` is an alternative implementation of mixed models specified via covariance structures:

```
library(mmrm)
e.mmrm <- mmrm(
  formula = FEV1 ~ RACE + SEX + ARMCD * AVISIT + us(AVISIT | USUBJID),
  data = fev_data
)
```

It leads nearly identical results compared to `lmm`:

```
e.lmm <- lmm(
  formula = FEV1 ~ RACE + SEX + ARMCD * AVISIT,
  repetition = ~ AVISIT | USUBJID, structure = "UN",
  data = fev_data, type.information = "expected"
)
```


Warning message:
In .lmmNormalizeData(as.data.frame(data)[unique(stats::na.omit(var.all))], :
    3 clusters have been removed.


```
logLik(e.mmrm) - logLik(e.lmm)
range(coef(e.mmrm) - coef(e.lmm))
range(vcov(e.mmrm) - vcov(e.lmm))
```

[1] -2.5413e-06
[1] -0.00018301 0.00016268
[1] -0.00039710 0.00020479

The main differences are:

- `mmrm` uses the expected information matrix to quantify uncertainty instead of the observed information matrix.

- `mmrm` implements the Kenward and Roger method for computing the degrees of freedom and not only the Satterthwaite approximation

- `mmrm` implements different covariance patterns

- `mmrm` is faster and probably more memory efficient

- `mmrm` has currently fewer post-processing methods (e.g. adjustment multiple comparisons when testing several model parameters). This being said, the latest version of the package (0.3.7) included several additional extractor of model feature so this may be improved in the future.

D.4   emmeans package

To illustrate a key difference between the emmeans package and the `effects.lmm` function we consider an informative and unbalanced group variable:

```
gastricbypassLB$group2 <- gastricbypassLB$weight1>150
```

Since `lmm`:

```
eCS.lmm_2 <- lmm(glucagonAUC ~ visit*group2, repetition =~visit|id, structure = "CS", data =
    gastricbypassLB)
logLik(eCS.lmm_2)
```

`[1] -315.2`

we will use the equivalent with the random effect specification:

```
eRI.lmer_2 <- lmer(glucagonAUC ~ visit*group2 + (1|id), data = gastricbypassLB)
logLik(eRI.lmer_2)
```

`'log Lik.' -315.2 (df=10)`

While the two models are equivalent, the average outcome output by `effects`:

```
effects(eCS.lmm_2, variable = NULL)
```

```
           Average counterfactual outcome

       estimate    se   df  lower  upper
(t=1)    32.317 4.426 64.3 23.476 41.158
(t=2)    29.653 4.535 65.2 20.598 38.709
(t=3)    77.308 4.535 65.1  68.25 86.366
(t=4)     51.95 4.426 64.3 43.109 60.791
```

substantially differ from the one of emmeans:

```
library(emmeans)
emmeans(eRI.lmer_2, specs=~visit)
```

NOTE: Results may be misleading due to involvement in interactions
 visit emmean   SE   df lower.CL upper.CL
 1   33.6 5.53 64.2   22.6   44.7
 2   32.0 5.57 64.4   20.9   43.2
 3   70.0 5.57 64.4   58.9   81.1
 4   47.2 5.53 64.2   36.1   58.2

Results are averaged over the levels of: group2
Degrees-of-freedom method: kenward-roger
Confidence level used: 0.95

This is because when averaging over the level of a covariate, emmeans considers *balanced groups*. In the example, the groups are not balanced:

```
table(gastricbypassLB$group2)/NROW(gastricbypassLB)
```


FALSE   TRUE
  0.8   0.2

Based on the group and timepoint specific means:

```
eCS.elmm_2 <- model.tables(effects(eCS.lmm_2, variable = "group2"))
eCS.elmm_2
```


  group2 visit estimate   se   df lower upper   p.value
1 FALSE   1   31.430 4.9484 64.349 21.545 41.314 2.4688e-08
2 FALSE   2   28.067 5.0996 65.383 17.884 38.251 6.6737e-07
3 FALSE   3   82.173 5.1008 65.211 71.986 92.359 0.0000e+00
4 FALSE   4   55.126 4.9484 64.349 45.241 65.010 0.0000e+00
5   TRUE   1   35.864 9.8967 64.349 16.095 55.633 5.7374e-04
6   TRUE   2   35.997 9.8967 64.349 16.228 55.766 5.4953e-04
7   TRUE   3   57.848 9.8967 64.349 38.079 77.617 1.8339e-07
8   TRUE   4   39.246 9.8967 64.349 19.477 59.015 1.8651e-04

We illustrate the difference:

- emmeans:

```
0.5*eCS.elmm_2[eCS.elmm_2$group2==FALSE,"estimate"]+0.5*eCS.elmm_2[eCS.elmm_2$group2==TRUE,"
    estimate"]
```


[1] 33.647 32.032 70.010 47.186

- effects:

```
0.8*eCS.elmm_2[eCS.elmm_2$group2==FALSE,"estimate"]+0.2*eCS.elmm_2[eCS.elmm_2$group2==TRUE,"
    estimate"]
```

[1] 32.317 29.653 77.308 51.950

The "emmeans" approach gives equal "weight" to the expected value of both group:

```
mu.group1 <-  as.double(coef(e.group)["(Intercept)"])
mu.group2 <-  as.double(coef(e.group)["(Intercept)"] + coef(e.group)["group2TRUE"])
p.group1 <- 14/20            ; p.group2 <- 6/20
c(emmeans = (mu.group1+mu.group2)/2, predict = mu.group1 * p.group1 + mu.group2 * p.group2)
```

 emmeans   predict
4.450435 4.514352

D.5 effectsize package (R^2 or η^2)

Partial η^2 can be computed based on `lmer` using the effectsize package:

```
library(effectsize)
eta_squared(eCS.lmer)
cat("\n")
```

```
# Effect Size for ANOVA (Type III)

Parameter   | Eta2 (partial) |       95% CI
-------------------------------------------
visit       |           0.64 | [0.50, 1.00]
group       |           0.01 | [0.00, 1.00]
visit:group |           0.19 | [0.03, 1.00]

- One-sided CIs: upper bound fixed at
```

and are approximately equal to what one can compute "manually":

```
eCS.Wald <- anova(eCS.lmm)$multivariate
eCS.Wald$df.num*eCS.Wald$statistic/(eCS.Wald$df.num*eCS.Wald$statistic+eCS.Wald$df.denom)
```

```
[1] 0.335374 0.033811 0.186290
```

The will not be true for heteroschedastic models:

```
eUN.Wald <- anova(eUN.lmm)$multivariate
eUN.Wald$df.num*eUN.Wald$statistic/(eUN.Wald$df.num*eUN.Wald$statistic+eUN.Wald$df.denom)
```

```
[1] 0.50787 0.17905 0.32380
```

compared to:

```
eta_squared(eUN.lmer)
cat("\n")
```

```
# Effect Size for ANOVA (Type III)

Parameter   | Eta2 (partial) |       95% CI
-------------------------------------------
visit       |           0.76 | [0.54, 1.00]
group       |           0.01 | [0.00, 1.00]
visit:group |           0.32 | [0.00, 1.00]

- One-sided CIs: upper bound fixed at
```

But in that case both may be misleading as the proportion of explained variance is not clearly defined.

D.6 MuMIn package (R^2)

```
library(MuMIn)
r.squaredGLMM(eCS.lmer)
cat("\n")
```

```
          R2m      R2c
[1,] 0.51728 0.62222
```

To reproduce these R2, we extract from the random intercept model:

- the residual variance

```
sigmaW <- sigma(eCS.lmm)[1,1]-sigma(eCS.lmm)[1,2]
```

- the variance of the random effect

```
sigmaB <- sigma(eCS.lmm)[1,2]
```

- the variance of the fitted values:

```
sigma2_XB <- var(fitted(eCS.lmm))
```

and evalutae the ratios:

```
c(R2m = sigma2_XB/(sigmaW + sigmaB + sigma2_XB),
  R2c = (sigma2_XB + sigmaB)/(sigmaW + sigmaB + sigma2_XB))
```

```
    R2m      R2c
0.52549 0.62865
```

D.7 stats package (partial residuals)

The function `residuals.lm` can be used to extract partial residuals from `lm` objects. For instance:

```
gastricbypassW$group <- as.factor(as.numeric(gastricbypassW$id)%%2)
eIID.lm <- lm(weight4 ~ group + weight1, data = gastricbypassW)
pRes.lm <- residuals(eIID.lm, type = "partial")
head(pRes.lm)
```

```
       group   weight1
1    7.19282    3.6648
2   -0.20504   31.7052
3    0.60631  -17.3352
4    6.44389   22.7052
5   -1.59403  -16.7352
6  -18.23382    8.4052
```

Those generally differ (by a constant) from the one provided by `residuals.lmm`:

```
eIID.lmm <- lmm(weight4 ~ group + weight1, data = gastricbypassW)
(residuals(eIID.lmm, type = "partial", variable = "group") - pRes.lm[,"group"])
(residuals(eIID.lmm, type = "partial", variable = "weight1") - pRes.lm[,"weight1"])
```

```
     1      2      3      4      5      6      7      8      9     10     11     12     13     14
2.0702 2.0702 2.0702 2.0702 2.0702 2.0702 2.0702 2.0702 2.0702 2.0702 2.0702 2.0702 2.0702 2.0702
    15     16     17     18     19     20
2.0702 2.0702 2.0702 2.0702 2.0702 2.0702
     1      2      3      4      5      6      7      8      9     10     11     12     13     14
106.22 106.22 106.22 106.22 106.22 106.22 106.22 106.22 106.22 106.22 106.22 106.22 106.22 106.22
    15     16     17     18     19     20
106.22 106.22 106.22 106.22 106.22 106.22
```

Indeed, `residuals.lm` centers the design matrix of the variable relative to which the partial residuals are computed:

```
coef(eIID.lm)["group1"] * mean(gastricbypassW$group=="1")
coef(eIID.lm)["weight1"] * mean(gastricbypassW$weight1)
```

```
group1
2.0702
weight1
 106.22
```

For continuous variable with a linear effect, these residuals can be obtained by setting the `type` argument to `"partial-center"`:

```
(residuals(eIID.lmm, type = "partial-center", variable = "weight1") - pRes.lm[,"weight1"])
```

```
          1           2           3           4           5           6           7           8
 1.7675e-13  6.7502e-14 -6.3949e-14  5.6843e-14 -3.9080e-14  8.1712e-14 -3.7303e-14  5.9508e-14
          9          10          11          12          13          14          15          16
-4.2633e-14  4.4409e-14 -2.9310e-14  5.5123e-14 -4.6185e-14  4.4409e-14 -4.2633e-14  4.6185e-14
         17          18          19          20
-3.9968e-14  5.3291e-14 -1.4211e-14  3.5527e-14
```

⚠ When evaluating the partial residuals relative to categorical variables, interactions, or non-linear terms, the output obtained with `partial-center` will not match the one of `residuals.lm`. Indeed `partial-center` will, when numeric, center the original variable whereas `residuals.lm` will center the column relative to the coefficient in the design matrix.